

Data Management for Data Science

Lecture 3: Principles of Data Management

Prof. Asoc. Endri Raço

Today's Lecture

1. Data Management
2. Data Models
3. RDBMs and the Relational Data Model

1. Data Management

Data Management

- Data represents the **traces** of real-world processes.
- Data is valuable **but** hard and costly to manage
 - Storage, representation complexity, collection
- Data management seeks to answer two questions:
 - What operations do we want to perform on this data?
 - What functionality do we need to manage this data?

Required Functionality

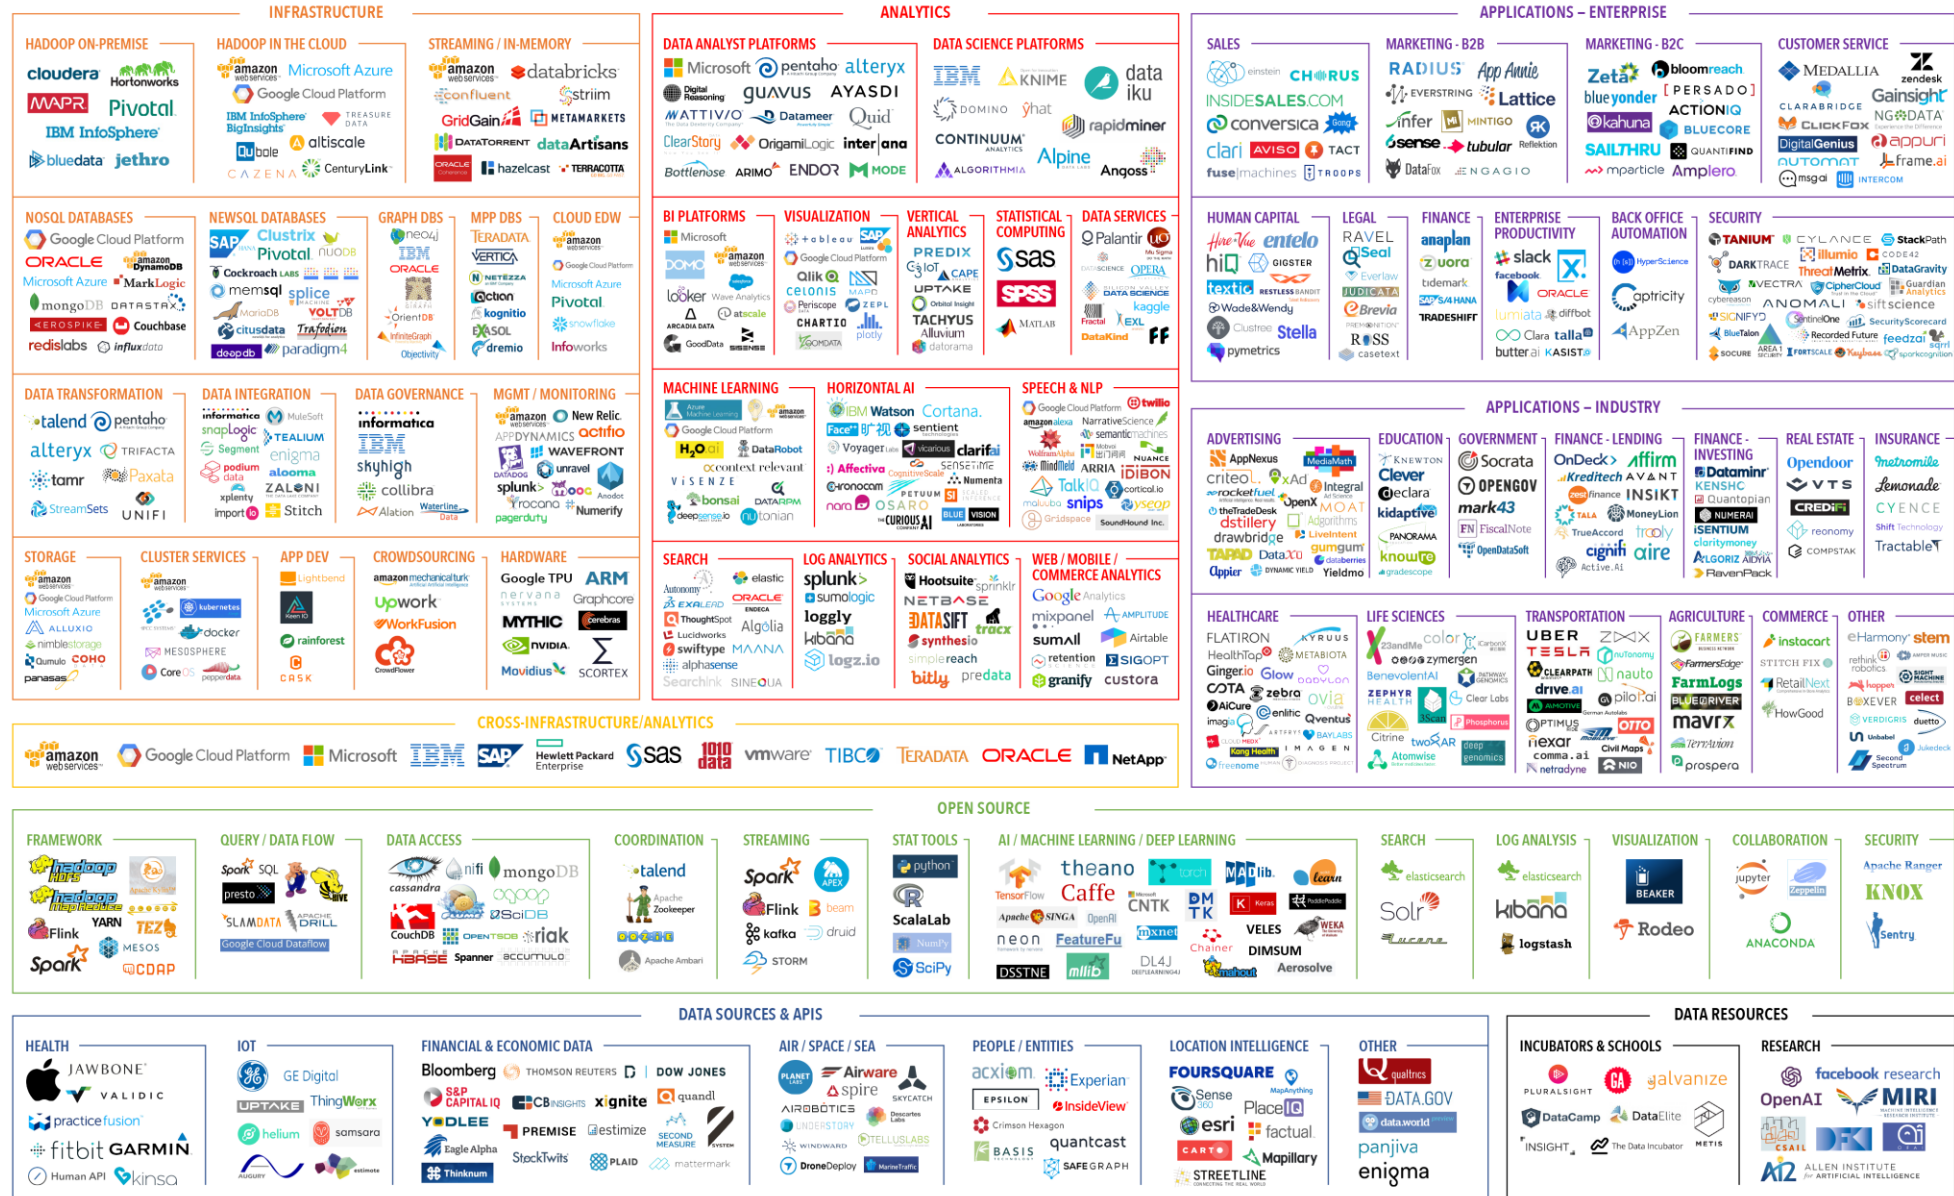
- Describe real-world entities in terms of stored data
- Create & persistently store large datasets
- Efficiently query & update
 - Must handle complex questions about the data
 - Must handle sophisticated updates
 - Performance matters
- Change structure (e.g., add attributes)
- Concurrency control: enable simultaneous queries, updates etc
- Crash recovery
- Access control, security, integrity

It is difficult and costly to implement all these features!

Systems providing data management features

- Relational database management systems
- HDFS-based systems (e.g., hadoop)
- Stream management systems: Apache Kafka
- Others?

BIG DATA LANDSCAPE 2017




2. Data Models

What you will learn about in this section

1. Types of Data
2. Data Models

Data is highly heterogeneous

- Structured data
- Semi-structured data
- Unstructured data



Increasing amounts of data

Structured data

- Information with a high degree of organization
- All data conforms to a schema. Ex: business data
- Easy to query, search over, aggregate
- Example: tables in a database, tables in excel, etc.

Semi-structured data

- Some structure in the data but implicit and irregular
- It contains tags or other markers to separate semantic elements and enforce hierarchies of records and fields within the data
- Example: JSON, HTML, XML

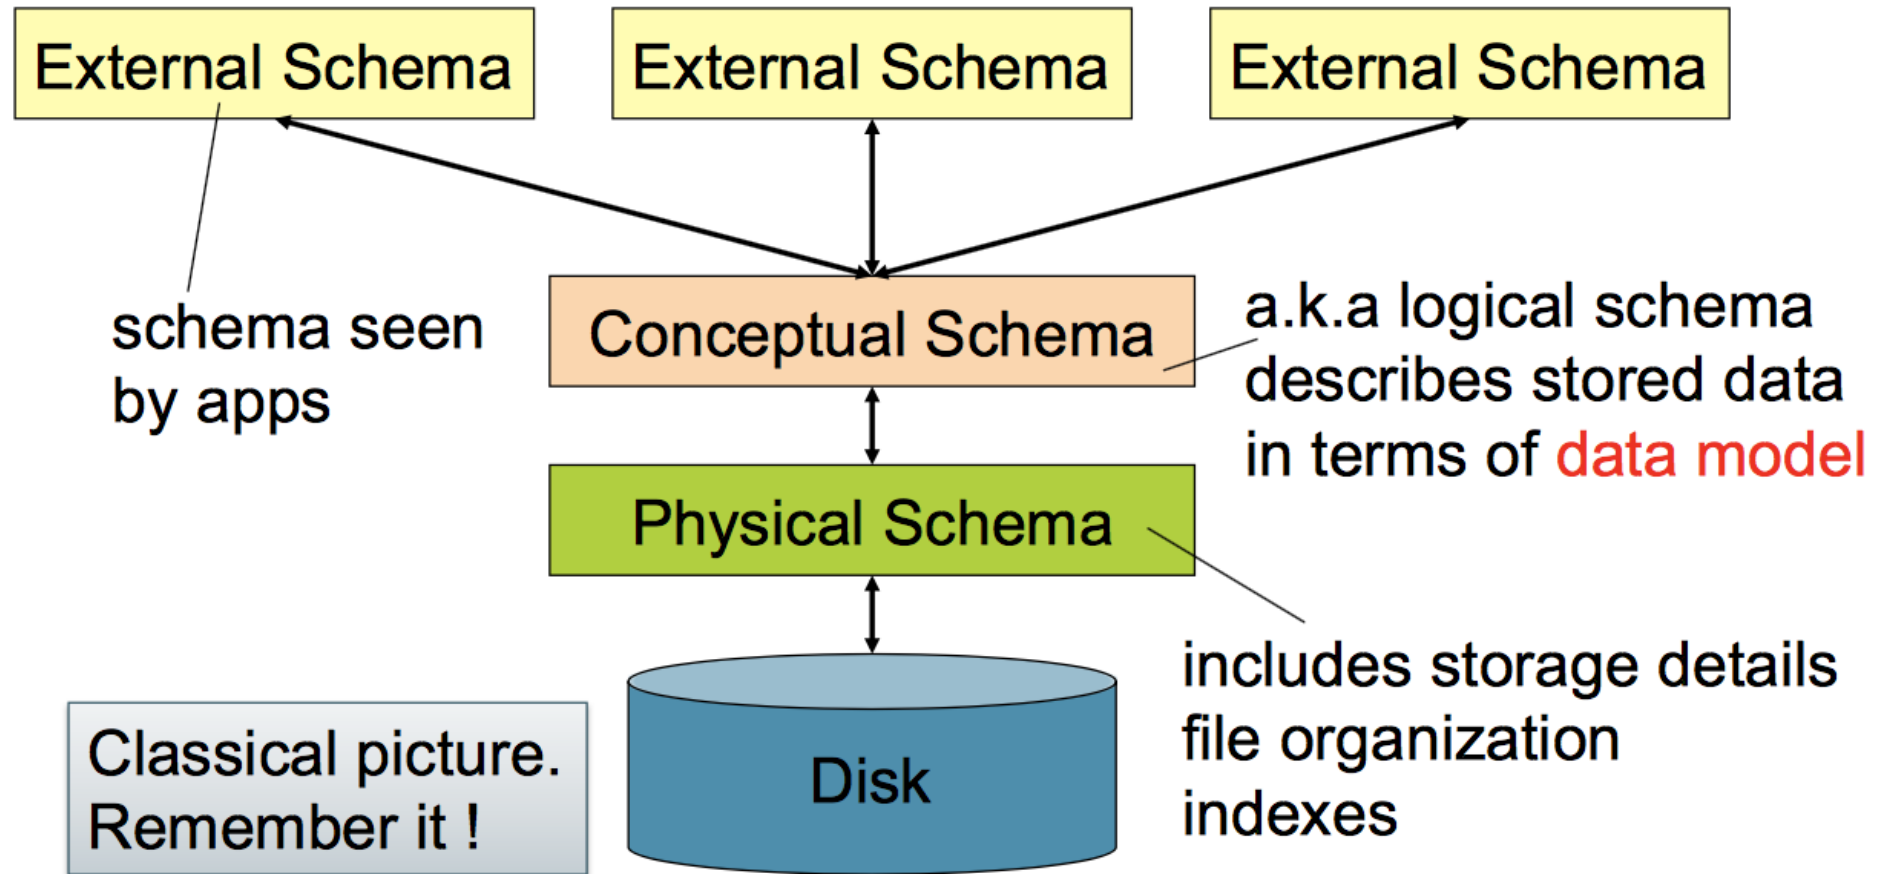
Unstructured data

- Information that either does not have a pre-defined structure or is not organized in a pre-defined manner.
- Text, video, images, etc.
- Abundant and extremely valuable. Hard to query, aggregate, analyze, search.

Data Model

- A **data model** is a collection of concepts for describing data
- A **schema** is a description of a particular collection of data, **using the given data model**
- A **data model** enables users to define the data using high-level constructs without worrying about many low-level details of how data will be stored on disk.

Levels of abstraction



Data models

- Relational **Most database management systems**
- Key/Value
- Graph
- Document
- Column-family
- Array/Matrix
- Hierarchical
- Network

Data models

- Relational
- Key/Value
- Graph
- Document
- Column-family
- Array/Matrix
- Hierarchical
- Network

No SQL

Data models

- Relational
- Key/Value
- Graph
- Document
- Column-family
- Array/Matrix
- Hierarchical
- Network

Machine learning, Scientific applications

Data models

- Relational
- Key/Value
- Graph
- Document
- Column-family
- Array/Matrix
- Hierarchical
- Network

Obsolete / Rare

3. RDBMs and the Relational Data Model

What you will learn about in this section

1. Definition of DBMS
2. Data models & the relational data model
3. Schemas & data independence

What is a DBMS?

- A large, integrated collection of data
- Models a real-world enterprise
 - *Entities* (e.g., Students, Courses)
 - *Relationships* (e.g., Alice is enrolled in CS564)

A Database Management System (DBMS) is a piece of software designed to store and manage databases

A Motivating, Running Example

- Consider building a course management system (**CMS**):

- Students
- Courses
- Professors

} *Entities*

- Who takes what
- Who teaches what

} *Relationships*

Data models

- A **data model** is a collection of concepts for describing data
 - The relational model of data is the most widely used model today
 - Main Concept: the *relation*- essentially, a table
- A **schema** is a description of a particular collection of data, **using the given data model**
 - E.g. every *relation* in a relational data model has a *schema* describing types, etc.

Modeling the Course Management System

- *Logical Schema*

- Students(sid: *string*, name: *string*, gpa: *float*)
- Courses(cid: *string*, cname: *string*, credits: *int*)
- Enrolled(sid: *string*, cid: *string*, grade: *string*)

sid	Name	Gpa
101	Bob	3.2
123	Mary	3.8

Students

Relations

sid	cid	Grade
123	564	A

Enrolled

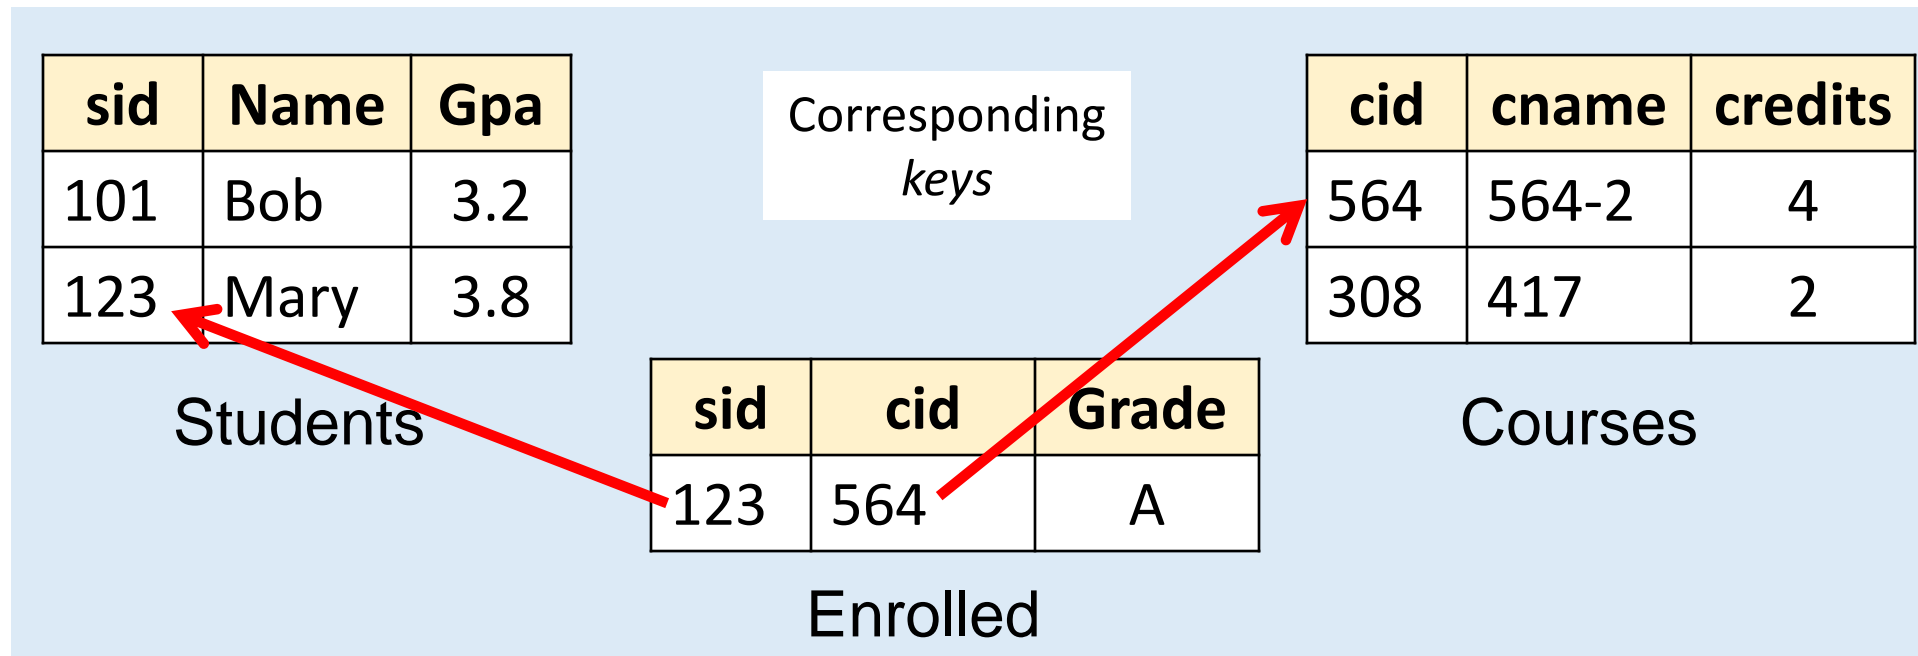
cid	cname	credits
564	564-2	4
308	417	2

Courses

Modeling the Course Management System

- *Logical Schema*

- Students(sid: *string*, name: *string*, gpa: *float*)
- Courses(cid: *string*, cname: *string*, credits: *int*)
- Enrolled(sid: *string*, cid: *string*, grade: *string*)



Other Schemata...

- *Physical Schema*: describes data layout
 - Relations as unordered files
 - Some data in sorted order (index)



Administrators

- *Logical Schema*: Previous slide



Applications

- *External Schema*: (Views)
 - Course_info(cid: *string*, enrollment: *integer*)
 - Derived from other tables

Data independence

Concept: Applications do not need to worry about *how the data is structured and stored*

Logical data independence:
protection from changes in the
logical structure of the data

I.e. should not need to ask: can we add a new entity or attribute without rewriting the application?

Physical data independence:
protection from *physical layout changes*

I.e. should not need to ask: which disks are the data stored on? Is the data indexed?

One of the most important reasons to use a DBMS

Relational Model

- **Structure:** The definition of relations and their contents.
- **Integrity:** Ensure the database's contents satisfy constraints.
- **Manipulation:** How to access and modify a database's contents.

Tables in the Relational Model

Product

PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

A relation or table is a multiset of tuples having the attributes specified by the schema

Let's break this definition down

Tables in the Relational Model

Product

PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

A multiset is an unordered list (or: a set with multiple duplicate instances allowed)

List: [1, 1, 2, 3]

Set: {1, 2, 3}

Multiset: {1, 1, 2, 3}

i.e. no *next()*, etc. methods!

Tables in the Relational Model

Product

PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

An attribute (or column) is a typed data entry present in each tuple in the relation

*Attributes must have an **atomic** type, i.e. not a list, set, etc.*

Tables in the Relational Model

Product

PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

*Also referred to sometimes as a **record***

A **tuple** or **row** is a single entry in the table having the attributes specified by the schema

Tables in the Relational Model

Product

PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

The number of tuples is the cardinality of the relation

The number of attributes is the arity of the relation

n -ary Relation
=
Table with n columns

Data Types in Relational Model

- Atomic types:
 - Characters: CHAR(20), VARCHAR(50)
 - Numbers: INT, BIGINT, SMALLINT, FLOAT
 - Others: MONEY, DATETIME, ...
- Every attribute must have an atomic type
 - Hence tables are flat

Table Schemas

- The **schema** of a table is the table name, its attributes, and their types:

```
Product(Pname: string, Price: float, Category: string, Manufacturer:  
string)
```

- A **key** is an attribute whose values are unique; we underline a key

```
Product(Pname: string, Price: float, Category: string, Manufacturer:  
string)
```

Key constraints

A key is a **minimal subset of attributes** that acts as a unique identifier for tuples in a relation

- A key is an implicit constraint on which tuples can be in the relation
 - i.e. if two tuples agree on the values of the key, then they must be the same tuple!

```
Students(sid:string, name:string, gpa: float)
```

1. Which would you select as a key?
2. Is a key always guaranteed to exist?
3. Can we have more than one key?

NULL and NOT NULL

- To say “don’t know the value” we use **NULL**
 - NULL has (sometimes painful) semantics, more details later

```
Students(sid:string, name:string, gpa: float)
```

sid	name	gpa
123	Bob	3.9
143	Jim	NULL

Say, Jim just enrolled in his first class.

We may constrain a column to be NOT NULL, e.g., “name” in this table

Foreign Key constraints

- A foreign key specifies that an attribute from one relation has to map to a tuple in another relation.

Foreign Key constraints

- Suppose we have the following schema:

Students(sid: string, name: string, gpa: float)

Enrolled(student_id: string, cid: string, grade: string)

- And we want to impose the following constraint:
 - 'Only real students may enroll in courses' i.e. a student must appear in the Students table to enroll in a class

Students

sid	name	gpa
101	Bob	3.2
123	Mary	3.8

Enrolled

student_id	cid	grade
123	564	A
123	537	A+

student_id alone is not a key- what is?

We say that student_id is a **foreign key** that refers to Students

Summary of Schema Information

- Schema and Constraints are how databases understand the semantics (meaning) of data
- They are also useful for optimization

DATA MANIPULATION LANGUAGES (DML)

- How to store and retrieve information from a database.
- Procedural: The query specifies the (high-level) strategy the DBMS should use to find the desired result.
- We will see SQL and Relational Algebra