

Data Management for Data Science

Lecture 8: Reasoning about Scale
& The MapReduce Abstraction

Prof.Asoc.Endri Raço

Today's Lecture

1. Scalability and Algorithmic Complexity
2. Data-Parallel Algorithms
3. The MapReduce Abstraction

1. Scalability and Algorithmic Complexity


What does scalable mean?

- Operationally:
 - Works even if the data does not fit in main memory
 - Use all available resources (cores/memory) on a single node (aka **scale up**)
 - Can make use of 1000s of cheap computers (cloud) – elastic (aka **scale out**)
- Algorithmically:
 - If you have N data items you should not perform more than N^m operations (polynomial complexity)
 - In many cases it should be $N \cdot \log(N)$ operations (streaming or too large data)
 - If you have N data items, you must do no more than N^m/k operations for some large k (k = number of cores/threads)


A sketch of algorithmic complexity

- Example: Find matching string sequences
- Given a set of string sequences
- Find all sequences equal to “GATTACGA”

Example: Find matching string sequences




Example: Find matching string sequences



Time = 0: TACCTGCC ? GATTACGA


Example: Find matching string sequences



Time = 0: TACCTGCC ? GATTACGA

No move cursor to next data entry


Example: Find matching string sequences



Time = 1: EFTAAGCA ? GATTACGA

No move cursor to next data entry


Example: Find matching string sequences



Time = 2: XXXXXXXX ? GATTACGA

No move cursor to next data entry


Example: Find matching string sequences



Time = n: GATTACGA ? GATTACGA


Yes! Output matching sequence

Example: Find matching string sequences



If we have 40 records we need to perform 40 comparisons


Example: Find matching string sequences




For N records we perform N comparisons

The algorithmic complexity is order N: $O(N)$

What if we knew the sequences are **sorted**




What if we knew the sequences are sorted



Time = 0: Start at 50% mark $\text{CTGTACA} < \text{GATTACGA}$


What if we knew the sequences are sorted



Time = 1: Start at 50% mark $CTGTACA < GATTACGA$

Skip to 75% mark (you know your sequence is in the second half)

What if we knew the sequences are sorted




Time = 2: We are at the 75% mark $TTGTCCA > GATTACGA$

Skip to 62.5% mark Match: $GATTACGA = GATTACGA$

We find our sequence in three steps. Now we can scan entries sequentially

What if we knew the sequences are **sorted**



How many comparisons?

For N records we did $\log(N)$ comparisons


The algorithm has complexity $O(\log(N))$ — much better scalability

2. Data-Parallel Algorithms

New task: Trim string sequences


- Given a set of string sequences
- Trim the final n characters of each sequence
- Generate a new dataset

New task: Trim string sequences (last 3 chars)




Time = 0: TACCTGCC -> TACCTG

New task: Trim string sequences (last 3 chars)



Time = 1: GATTCTGC -> GATTC


New task: Trim string sequences (last 3 chars)



Time = 2: CCCGAAT -> CCCG

Can we use a data structure to speed this operation?

New task: Trim string sequences (last 3 chars)




Time = 2: CCCGAAT -> CCCG


Can we use a data structure to speed this operation?

No. We have to touch every record! The task is $O(N)$.


New task: Trim string sequences (last 3 chars)



New task: Trim string sequences (last 3 chars)




New task: Trim string sequences (last 3 chars)




Time = 1: Process first element of each group

New task: Trim string sequences (last 3 chars)



Time = 2: Process second element of each group


New task: Trim string sequences (last 3 chars)



Time = 3: Process third element of each group


Etc.. How much time does this take?

New task: Trim string sequences (last 3 chars)








We only need $O(N/k)$ operations where k is the number of groups (workers)

Schematic of Parallel Algorithms







Schematic of Parallel Algorithms

2. You distribute the reads among k computers (workers)



3. Apply function f to each read (for every item in each chunk)



4. Obtain a big distributed set of outputs

Applications of parallel algorithms

- Convert TIFF images to PNG
- Run thousands of simulations for different model parameters
- Find the most common word in each document
- Compute the word frequency of every word in a single document
- Etc....

Applications of parallel algorithms

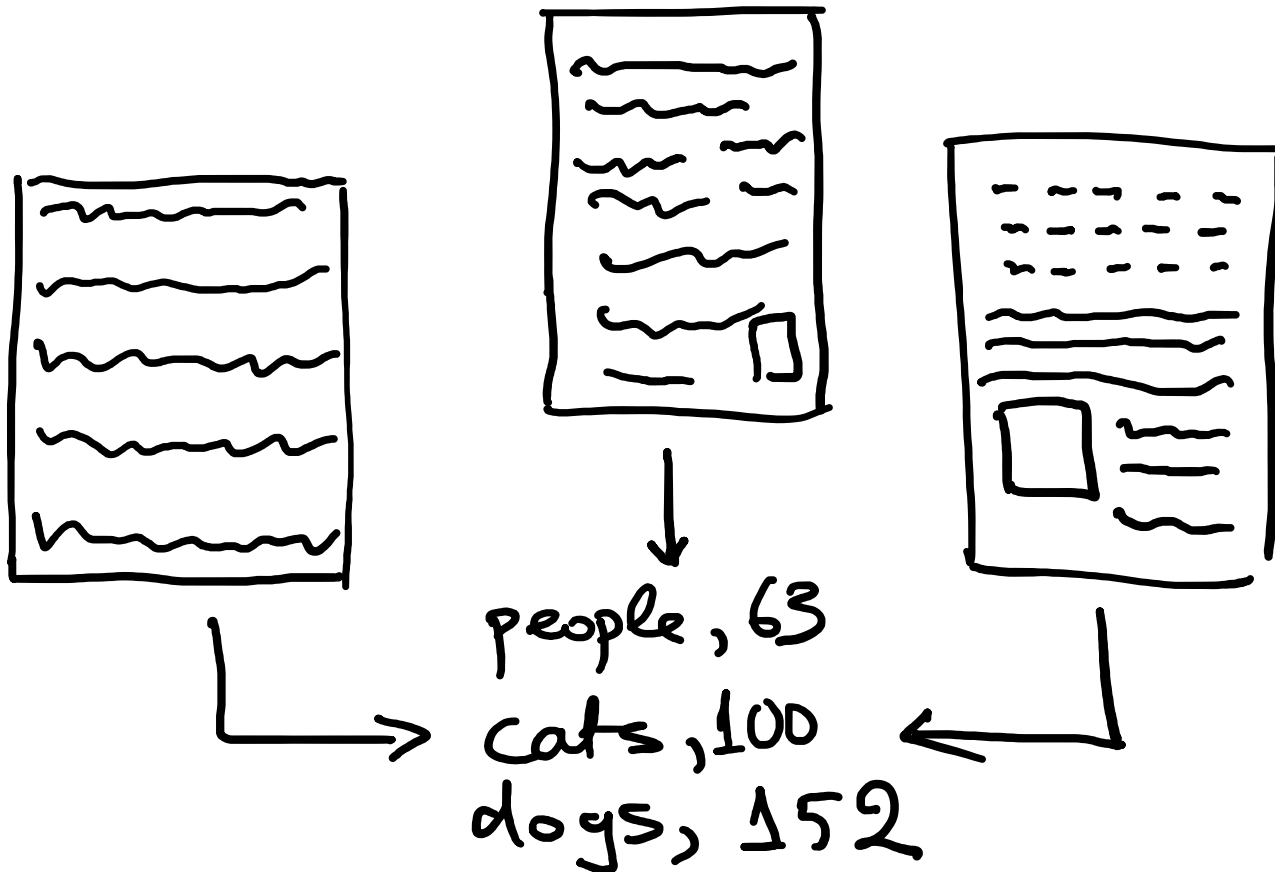
- Convert TIFF images to PNG
 - Run thousands of simulations for different model parameters
 - Find the most common word in each document
 - Compute the word frequency of every word in a single document
 - Etc....
-
- There is a common pattern in all these applications

Applications of parallel algorithms

- A function that ***maps*** a string to a trimmed string
- A function that ***maps*** a TIFF images to a PNG image
- A function that ***maps*** a set of parameters to simulation results
- A function that ***maps*** a document to its most common word
- A function that ***maps*** a document to a histogram of word frequencies

Applications of parallel algorithms

- What if we want to compute the word frequency across *all* documents?







3. The MapReduce Abstraction

Compute the word frequency across 5M documents

Millions of Documents



Distribute among k workers



for each doc return
(word, freq) pairs ↓

map

map
↓

map
↓

map
↓




map
↓

Compute the word frequency across 5M documents

Millions of Documents



Distribute among k workers



for each doc return (word, freq) pairs ↓

map ↓

map ↓

map ↓

map ↓


map ↓

Then what?



Challenge: in this task

- How can we make sure that a single computer has access to every occurrence of a given word regardless of which document it appeared in?
- *Ideas?*



Compute the word frequency **across** 5M documents




Compute the word frequency **across** 5M documents



Compute the word frequency **across** 5M documents



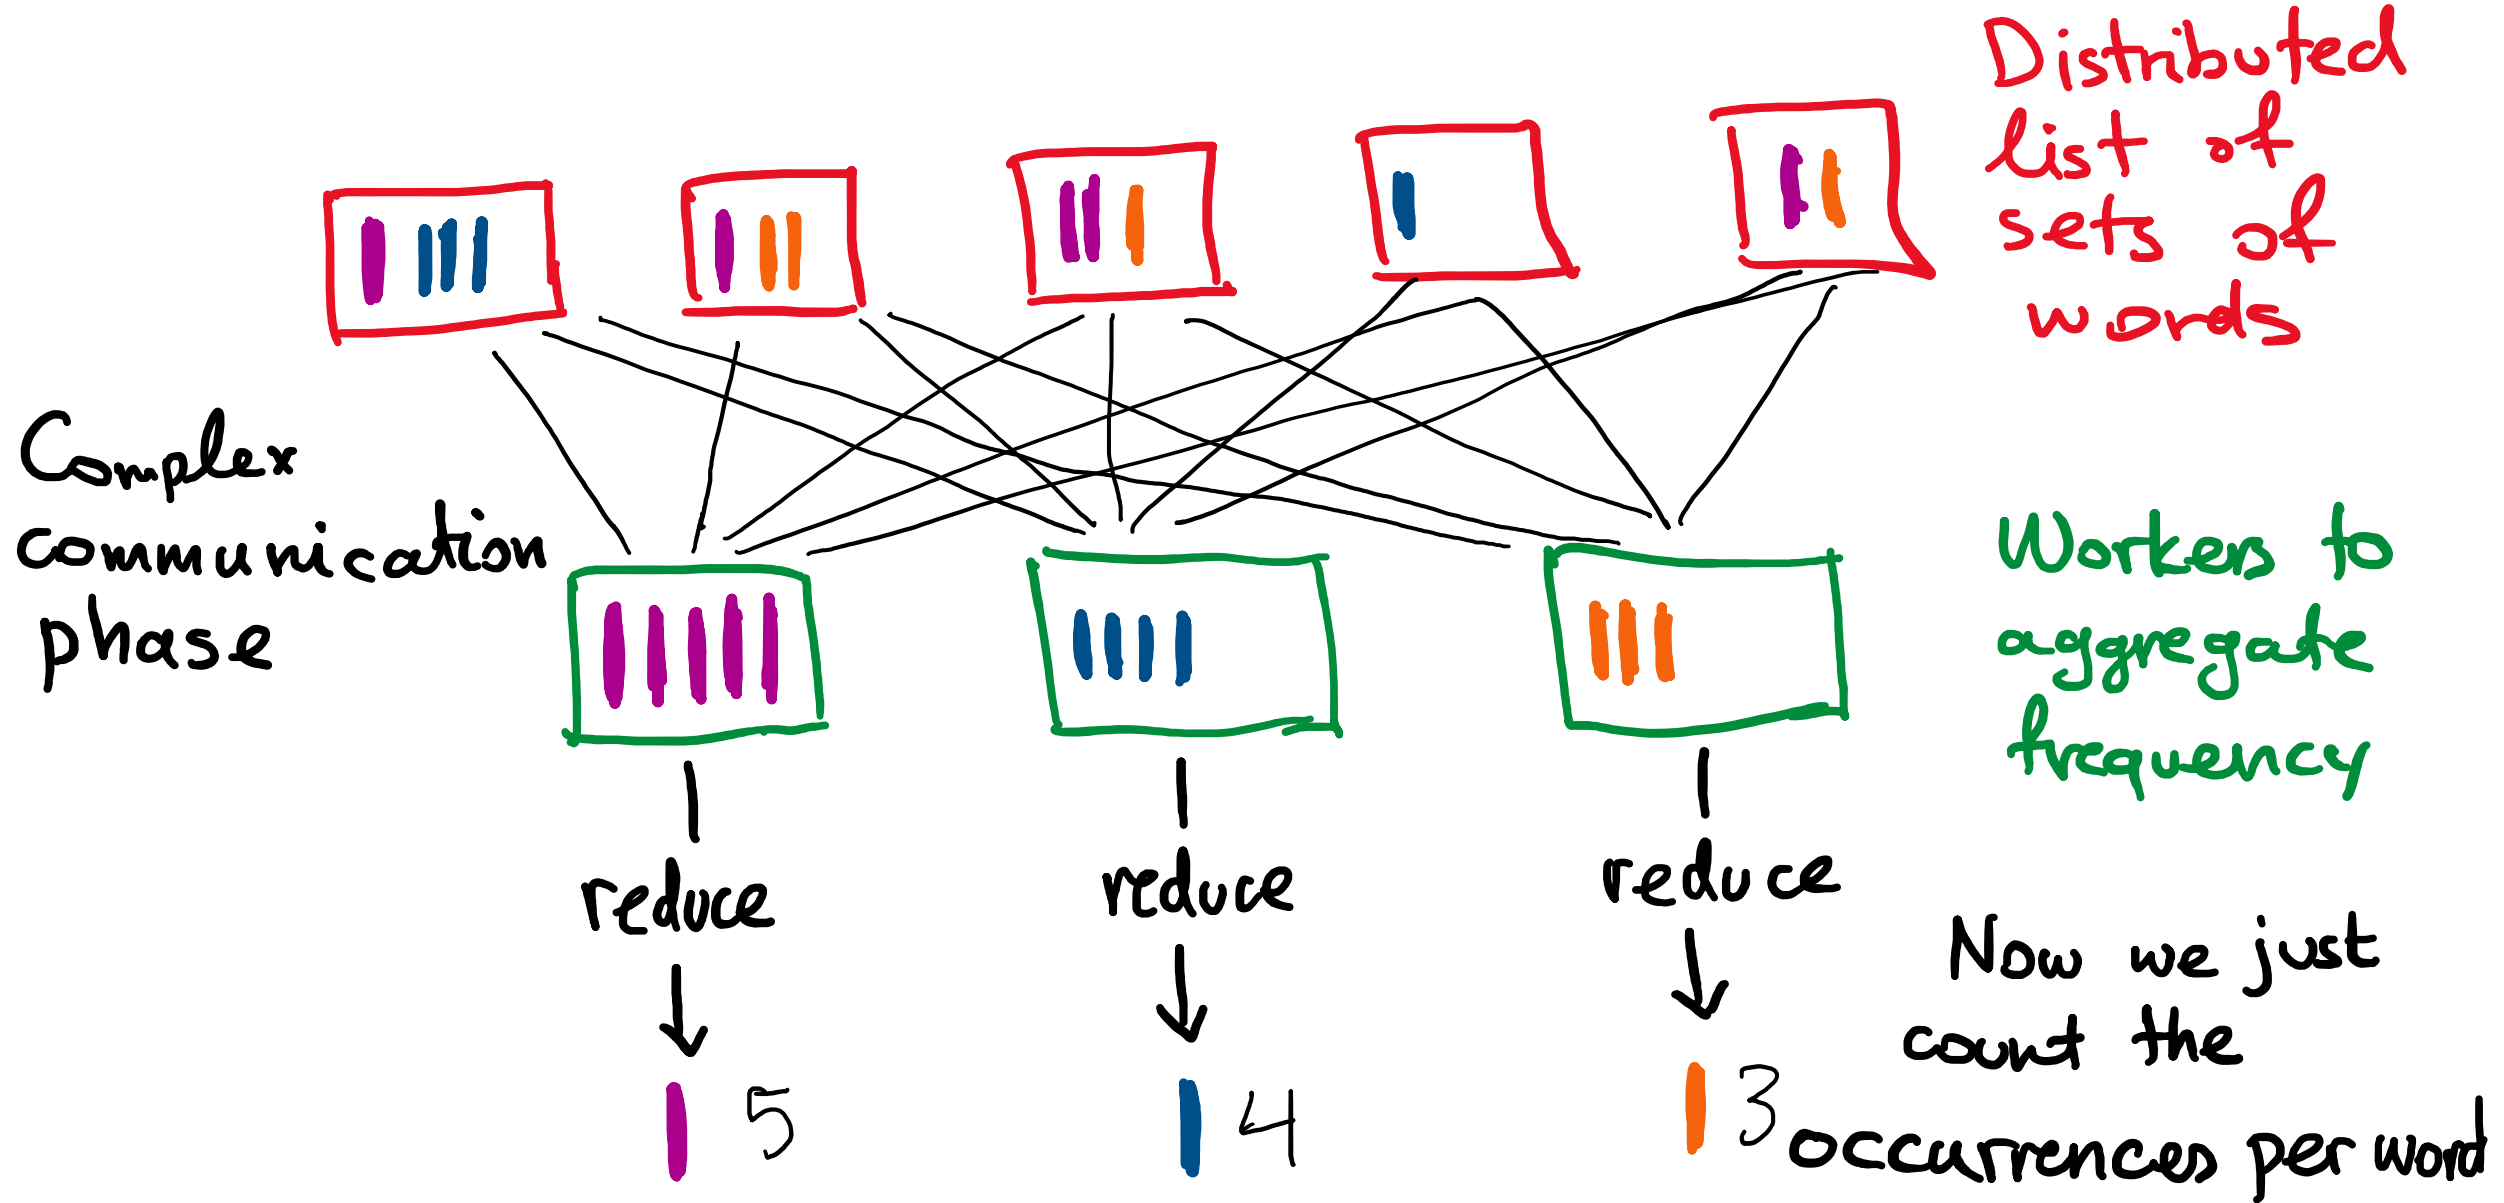
Compute the word frequency **across** 5M documents



We use a hash function here!

A hash function is any function that can be used to map data of arbitrary size to a data of a fixed size

Compute the word frequency across 5M documents



The Map Reduce Abstraction for Distributed Algorithms

