

Data Management for Data Science

Lecture 11: Spark

Prof. Asoc. Endri Raço

Logistics/Announcements


- Questions on PA3?

Today's Lecture

1. MapReduce Implementation
2. Spark

1. MapReduce Implementation

Recall: The Map Reduce Abstraction for Distributed Algorithms



MapReduce: what happens in between?

- **Map**

- Grab the relevant data from the source (parse into key, value)
- Write it to an intermediate file

- **Partition**

- Partitioning: identify which of R reducers will handle which keys
- Map partitions data to target it to one of R Reduce workers based on a partitioning function (both R and partitioning function user defined)

Map Worker

- **Shuffle & Sort**

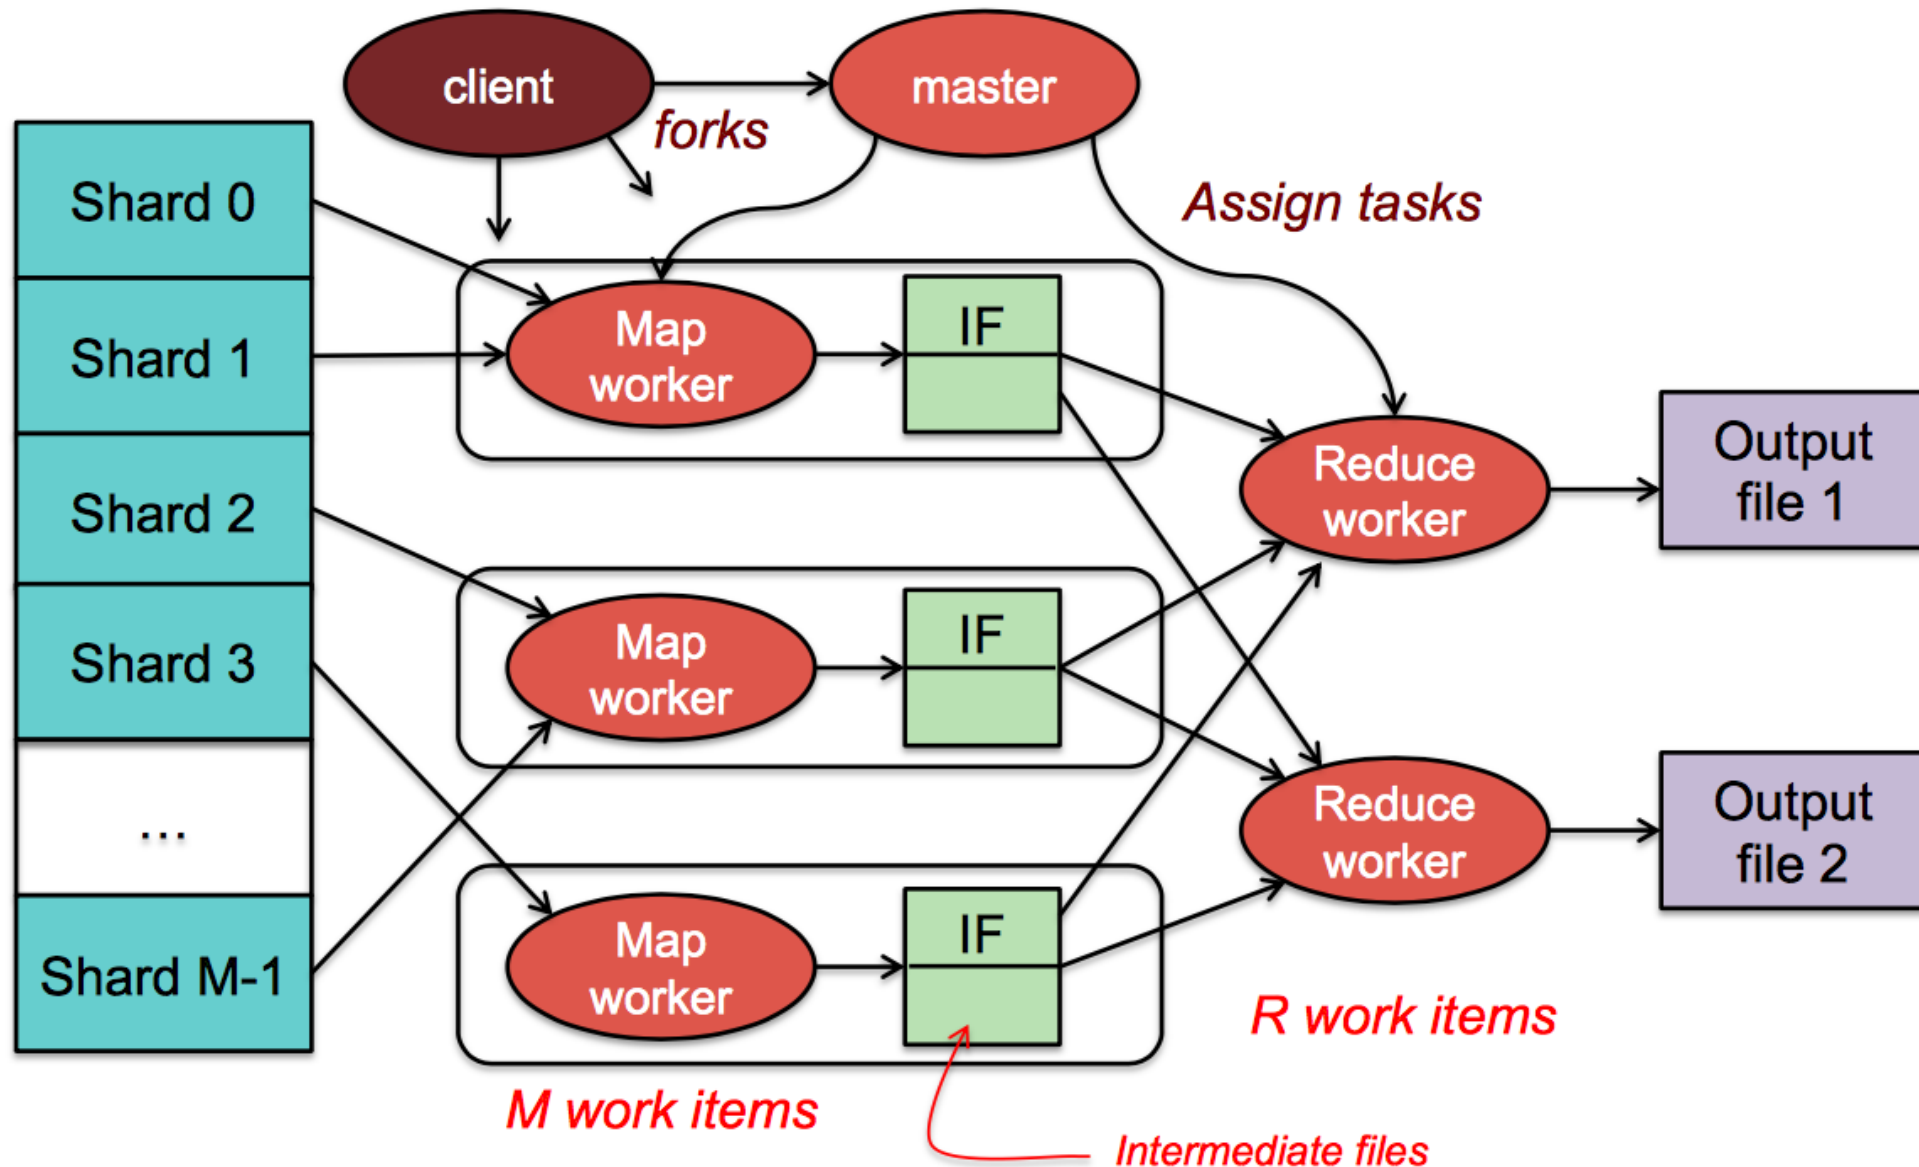
- Shuffle: Fetch the relevant partition of the output from all mappers
- Sort by keys (different mappers may have sent data with the same key)

- **Reduce**

- Input is the sorted output of mappers
- Call the user *Reduce* function per key with the list of values for that key to aggregate the results

Reduce Worker

MapReduce: the complete picture



Step 1: Split input files into chunks (shards)

- Break up the input data into M pieces (typically 64 MB)




Input files

Divided into M shards


Step 2: Fork processes

- Start up many copies of the program on a cluster of machines
 - **One master**: scheduler & coordinator
 - Lots of workers
- Idle workers are assigned either:
 - **map tasks** (each works on a shard) – there are M map tasks
 - **reduce tasks** (each works on intermediate files) – there are R
 - $R = \#$ partitions, defined by the user




Step 3: Run Map Tasks

- Reads contents of the input shard assigned to it
- Parses key/value pairs out of the input data
- Passes each pair to a user-defined *map* function
 - Produces intermediate key/value pairs
 - These are buffered in memory



Step 4: Create intermediate files

- Intermediate key/value pairs produced by the user's *map* function buffered in memory and are periodically written to the local disk
 - Partitioned into R regions by a **partitioning function**

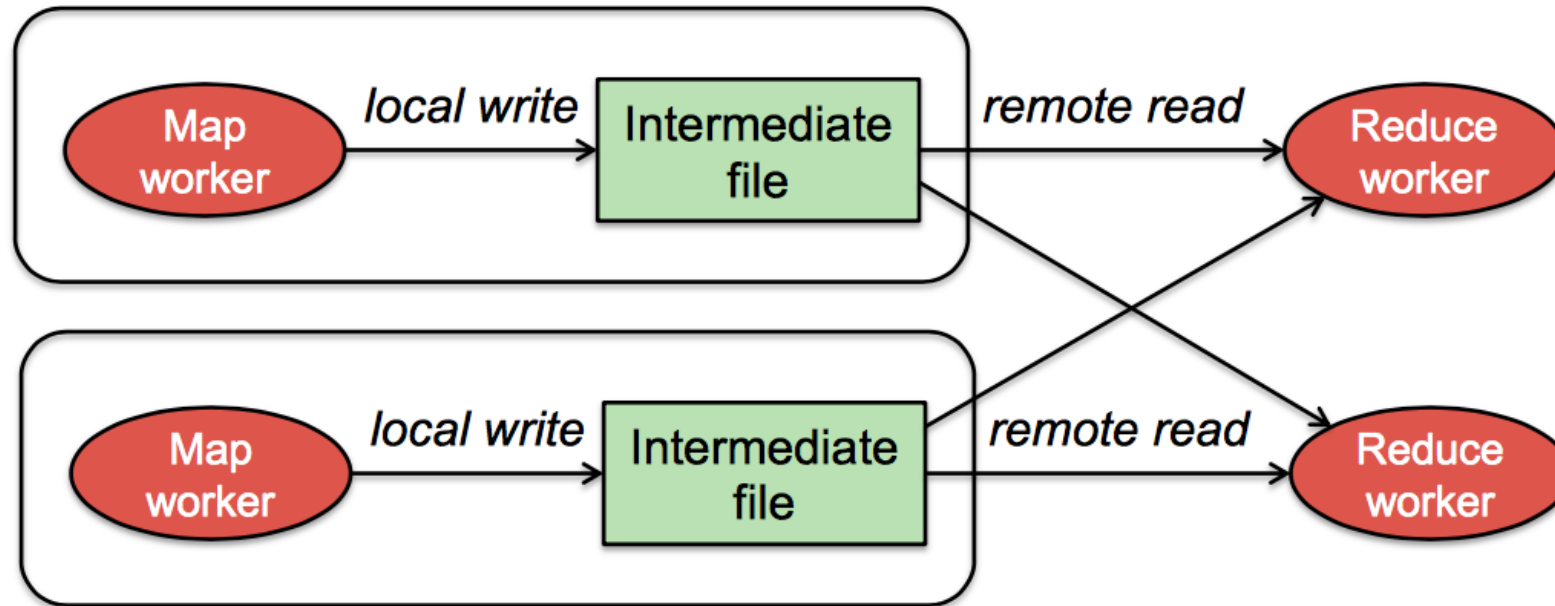


Step 4a: Partitioning

- Map data will be processed by Reduce workers
 - User's *Reduce* function will be called once per unique key generated by *Map*.
- We first need to **sort** all the (*key*, *value*) data by keys and decide which Reduce worker processes which keys
 - The Reduce worker will do the sorting
- **Partition function**
Decides which of R reduce workers will work on which key
 - Default function: $\text{hash}(\text{key}) \bmod R$
 - Map worker partitions the data by keys
- Each Reduce worker will later read their partition from every Map worker

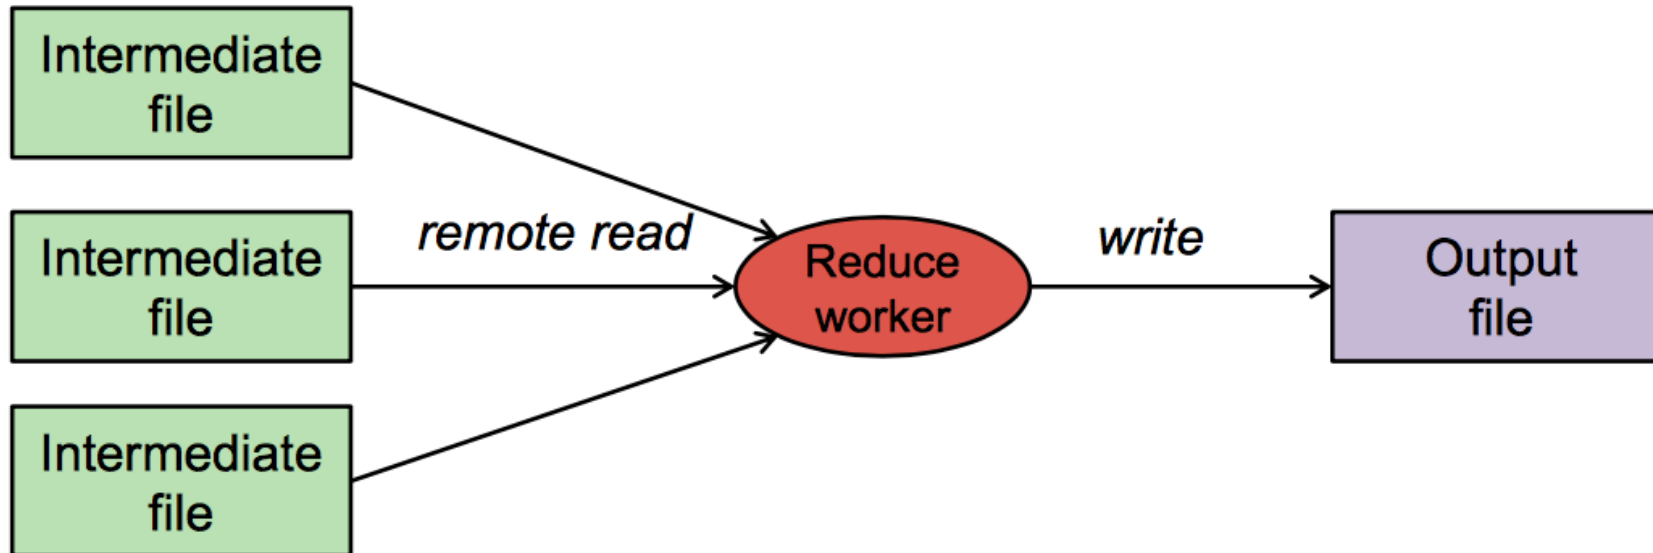
Step 5: Reduce Task - sorting

- Reduce worker gets notified by the master about the location of intermediate files for its partition
- **Shuffle**: Uses RPCs to read the data from the local disks of the map workers
- **Sort**: When the *reduce* worker reads intermediate data for its partition
 - It sorts the data by the intermediate keys
 - All occurrences of the same key are grouped together



Step 6: Reduce Task - reduce


- The sort phase grouped data with a unique intermediate key
- User's **Reduce** function is given the key and the set of intermediate values for that key
< key, (value1, value2, value3, value4, ...) >
- The output of the *Reduce* function is appended to an output file



Step 7: Return to user

- When all *map* and *reduce* tasks have completed, the master wakes up the user program
- The *MapReduce* call in the user program returns and the program can resume execution.
 - Output of *MapReduce* is available in *R* output files

MapReduce: the complete picture



2. Spark


Intro to Spark

- Spark is really a different implementation of the MapReduce programming model
- What makes Spark different is that it operates on Main Memory
- Spark: we write programs in terms of operations on resilient distributed datasets (RDDs).
- RDD (simple view): a collection of elements partitioned across the nodes of a cluster that can be operated on in parallel.
- RDD (complex view): RDD is an interface for data transformation, RDD refers to the data stored either in persisted store (HDFS) or in cache (memory, memory+disk, disk only) or in another RDD

RDDs in Spark

RDD: Resilient Distributed Datasets

- **Like a big list:**
 - Collections of objects spread across a cluster, stored in RAM or on Disk
- **Built through parallel transformations**
- **Automatically rebuilt on failure**



Operations


- **Transformations** (e.g. map, filter, groupBy)
- **Make sure input/output match**

MapReduce vs Spark

<satish, 26000>	<gopal, 50000>	<satish, 26000>	<satish, 26000>
<Krishna, 25000>	<Krishna, 25000>	<kiran, 45000>	<Krishna, 25000>
<Satishk, 15000>	<Satishk, 15000>	<Satishk, 15000>	<manisha, 45000>
<Raju, 10000>	<Raju, 10000>	<Raju, 10000>	<Raju, 10000>



Map and reduce
tasks operate on key-value
pairs




Spark operates on **RDD**

RDDs

- Partitions are recomputed on failure or cache eviction
- Metadata stored for interface:
 - Partitions – set of data splits associated with this RDD
 - Dependencies – list of parent RDDs involved in computation
 - Compute – function to compute partition of the RDD given the parent partitions from the Dependencies
 - Preferred Locations – where is the best place to put computations on this partition (data locality)
 - Partitioner – how the data is split into partitions

RDDs

Lazy computations model



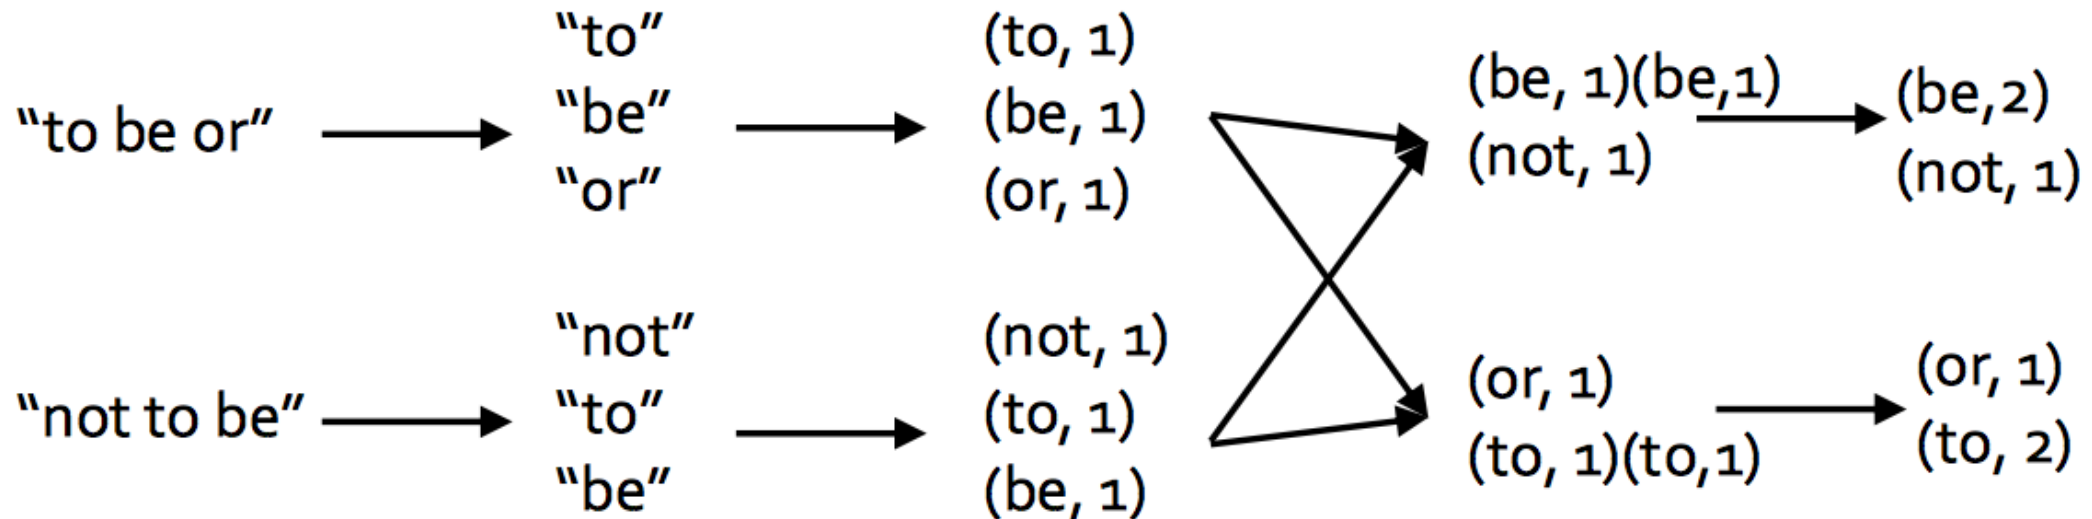
Transformation cause only metadata change

DAG


- Directed Acyclic Graph – sequence of computations performed on data
- Node – RDD partition
- Edge – transformation on top of the data
- Acyclic – graph cannot return to the older partition
- Directed – transformation is an action that transitions data partitions state (from A to B)

Example: Word Count


```
> lines = sc.textFile("hamlet.txt")  
> counts = lines.flatMap(lambda line: line.split(" "))  
                   .map(lambda word: (word, 1))  
                   .reduceByKey(lambda x, y: x + y)
```



Spark Architecture



Spark Components



Spark Driver


- Entry point of the Spark Shell (Scala, Python, R)
- The place where SparkContext is created
- Translates RDD into the execution graph
- Splits graph into stages
- Schedules tasks and controls their execution
- Stores metadata about all the RDDs and their partitions
- Brings up Spark WebUI with job information

Spark Executor

- Stores the data in cache in JVM heap or on HDDs
- Reads data from external sources
- Writes data to external sources
- Performs all the data processing

Dag Scheduler

- **General task graphs**
- **Automatically pipelines functions**
- **Data locality aware**
- **Partitioning aware to avoid shuffles**




More RDD Operations

- | | | |
|-------------------------|----------------------|-------------|
| • map | • reduce | sample |
| • filter | • count | take |
| • groupBy | • fold | first |
| • sort | • reduceByKey | partitionBy |
| • union | • groupByKey | mapWith |
| • join | • cogroup | pipe |
| • leftOuterJoin | • cross | save ... |
| • rightOuterJoin | • zip | |

Spark's secret is really the RDD abstraction

RDD: Resilient Distributed Datasets


- **Like a big list:**
 - Collections of objects spread across a cluster, stored in RAM or on Disk
- **Built through parallel transformations**
- **Automatically rebuilt on failure**



Operations

- **Transformations** (e.g. map, filter, groupBy)
- **Make sure input/output match**

Typical NoSQL architecture




CAP theorem for NoSQL

- What the CAP theorem really says: If you cannot limit the number of faults and requests can be directed to any server and you insist on serving every request you receive then you cannot possibly be consistent
- How it is interpreted: You must always give something up: consistency, availability or tolerance to failure and reconfiguration

CAP theorem for NoSQL

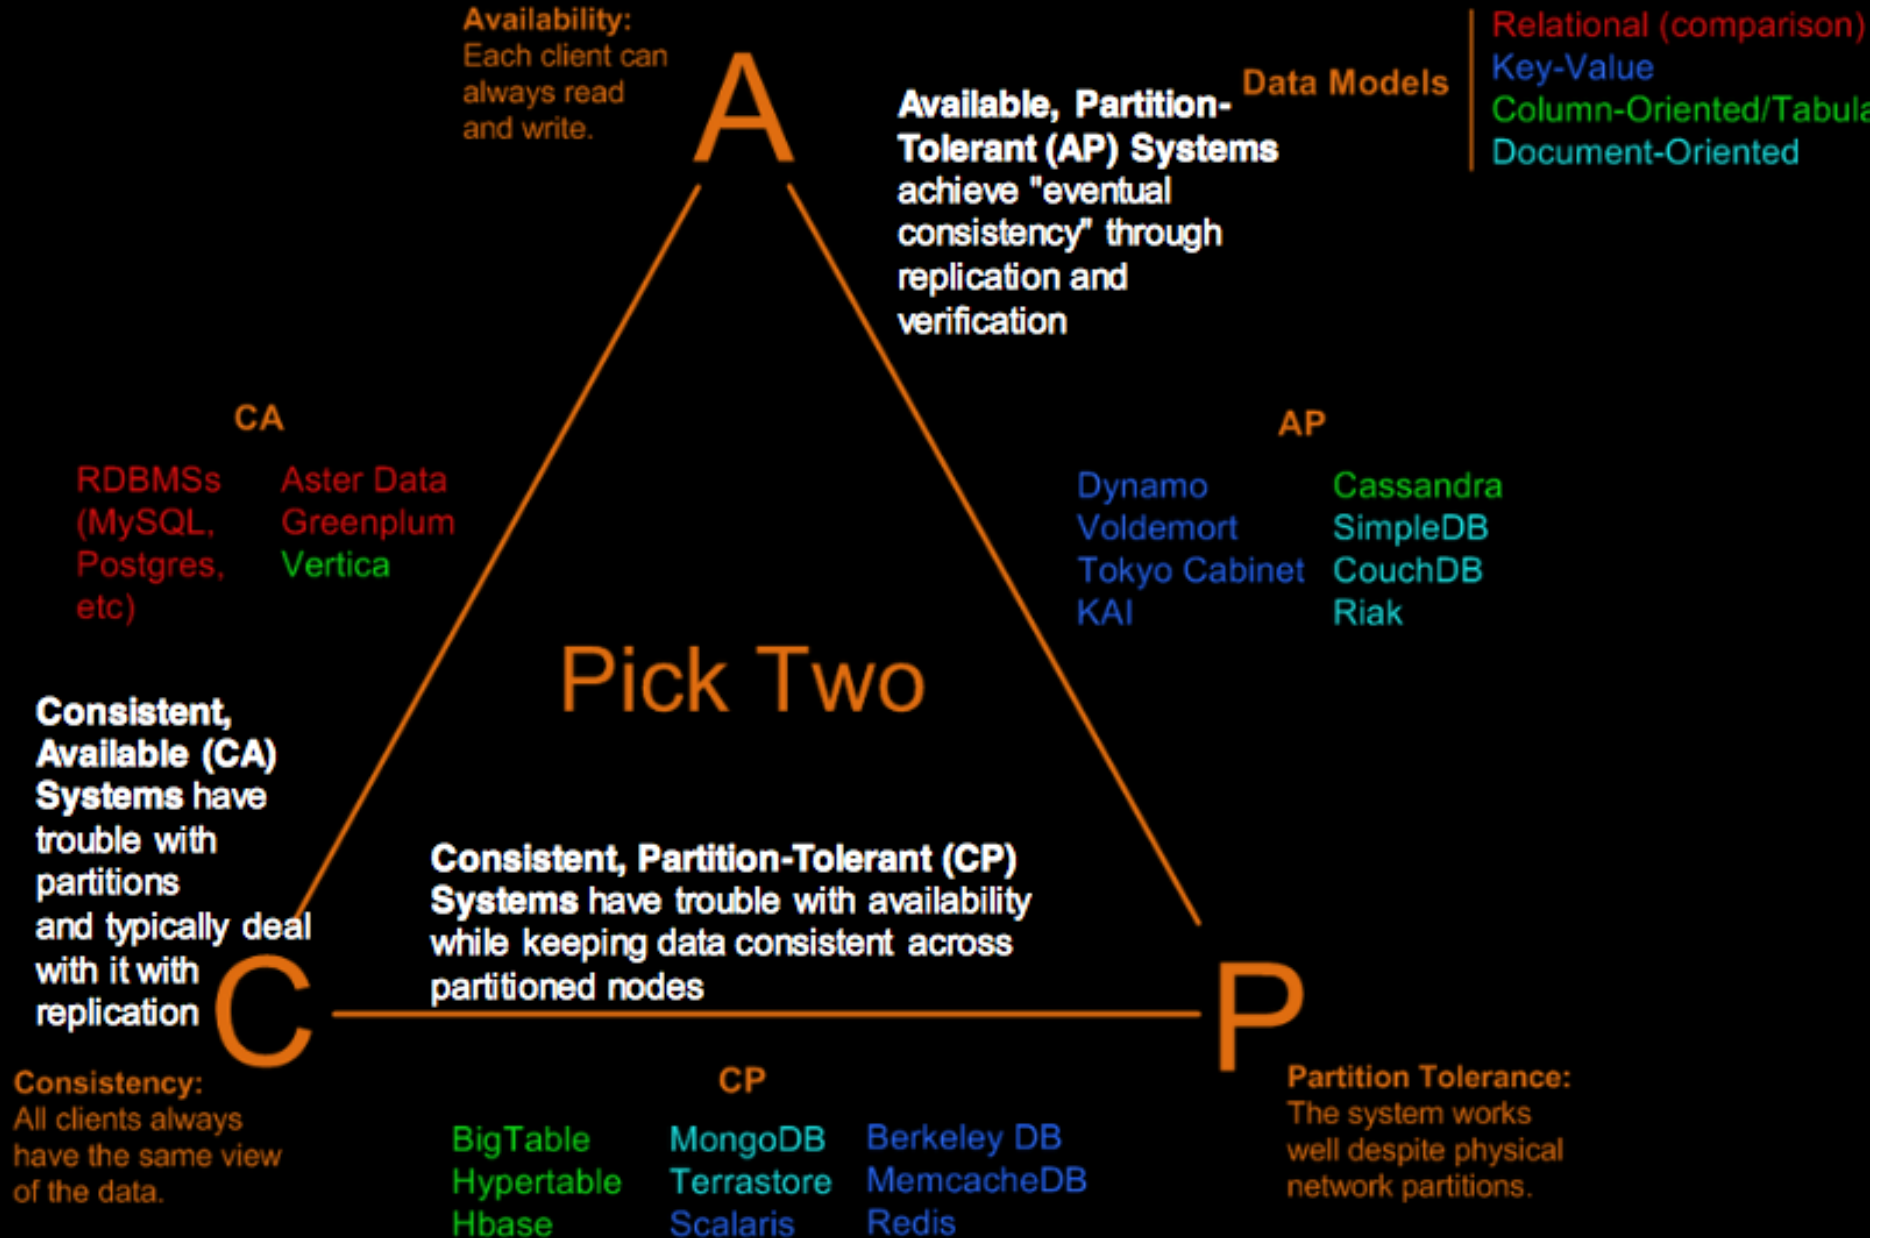
GIVEN:

- Many nodes
- Nodes contain **replicas of partitions** of the data
- **C**onsistency
 - All replicas contain the same version of data
 - Client always has the same view of the data (no matter what node)
- **A**vailability
 - System remains operational on failing nodes
 - All clients can always read and write
- **P**artition tolerance
 - multiple entry points
 - System remains operational on system split (communication malfunction)
 - System works well across physical network partitions



CAP Theorem:
satisfying all three at the
same time is impossible

Visual Guide to NoSQL Systems




Sharding of data

- Distributes a single logical database system across a cluster of machines
- Uses range-based partitioning to distribute documents based on a specific shard key
- Automatically balances the data associated with each shard
- Can be turned on and off per collection (table)

Replica Sets

- Redundancy and Failover
- Zero downtime for upgrades and maintenance
- Master-slave replication
 - Strong Consistency
 - Delayed Consistency
- Geospatial features



How does NoSQL vary from RDBMS?

- Looser schema definition
- Applications written to deal with specific documents/ data
 - Applications aware of the schema definition as opposed to the data
- Designed to handle distributed, large databases
- Trade offs:
 - No strong support for ad hoc queries but designed for speed and growth of database
 - Query language through the API
 - Relaxation of the ACID properties

Benefits of NoSQL

Elastic Scaling

- RDBMS scale up – bigger load , bigger server
- NO SQL scale out – distribute data across multiple hosts seamlessly

DBA Specialists

- RDMS require highly trained expert to monitor DB
- NoSQL require less management, automatic repair and simpler data models

Big Data

- Huge increase in data
RDMS: capacity and constraints of data volumes at its limits
- NoSQL designed for big data

Benefits of NoSQL

Flexible data models

- Change management to schema for RDMS have to be carefully managed
- NoSQL databases more relaxed in structure of data
 - Database schema changes do not have to be managed as one complicated change unit
 - Application already written to address an amorphous schema

Economics

- RDMS rely on expensive proprietary servers to manage data
- No SQL: clusters of cheap commodity servers to manage the data and transaction volumes
- Cost per gigabyte or transaction/second for NoSQL can be lower than the cost for a RDBMS

Drawbacks of NoSQL

- Support

- RDBMS vendors provide a high level of support to clients
 - Stellar reputation
- NoSQL – are open source projects with startups supporting them
 - Reputation not yet established

- Maturity

- RDMS mature product: means stable and dependable
 - Also means old no longer cutting edge nor interesting
- NoSQL are still implementing their basic feature set

Drawbacks of NoSQL

- **Administration**

- RDMS administrator well defined role
- No SQL's goal: no administrator necessary however NO SQL still requires effort to maintain

- **Lack of Expertise**

- Whole workforce of trained and seasoned RDMS developers
- Still recruiting developers to the NoSQL camp

- **Analytics and Business Intelligence**

- RDMS designed to address this niche
- NoSQL designed to meet the needs of an Web 2.0 application - not designed for ad hoc query of the data
 - Tools are being developed to address this need

ACID or BASE

