

# Data Management for Data Science

Lecture 16: Intro to ML and Decision Trees

Prof. Asoc. Endri Raço

# Today's Lecture

1. Intro to Machine Learning
2. Types of Machine Learning
3. Decision Trees

# 1. Intro to Machine Learning

# What is Machine Learning?

- “Learning is any process by which a system improves performance from experience” – Herbert Simon

- Definition by Tom Mitchell (1998):

Machine Learning is the study of algorithms that

- Improve their performance  $P$
- at some task  $T$
- with experience  $E$

A well-defined learning task is given by  $\langle P, T, E \rangle$ .

# What is Machine Learning?

Machine Learning is the study of algorithms that

- Improve their performance  $P$
- at some task  $T$
- with experience  $E$

A well-defined learning task is given by  $\langle P, T, E \rangle$ .


**Experience:** data-driven task, thus statistics, probability

**Example:** use height and weight to predict gender


# When do we use machine learning?

ML is used when:

- Human expertise does not exist (navigating on Mars)
- Humans can't explain their expertise (speech recognition)
- Models must be customized (personalized medicine)
- Models are based on huge amounts of data (genomics)




# A task that requires machine learning



What makes a hand drawing be 2?

# Modern machine learning: Autonomous cars





# Modern machine learning: Scene Labeling



# Modern machine learning: Speech Recognition




## 2. Types of Machine Learning

# Types of Learning

- Supervised (inductive) learning
  - Given: training data + desired outputs (labels)
- Unsupervised learning
  - Given: training data (without desired outputs)
- Semi-supervised learning
  - Given: training data + a few desired outputs
- Reinforcement learning
  - Rewards from sequence of actions


# Supervised Learning: Regression

- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$ 
  - $y$  is real-valued == regression




# Supervised Learning: Classification

- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$ 
  - $y$  is categorical == regression



# Supervised Learning: Classification


- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$ 
  - $y$  is categorical == regression





# Supervised Learning

- Value  $x$  can be multi-dimensional.
  - Each dimension corresponds to an attribute




- Clump Thickness
- Uniformity of Cell Size
- Uniformity of Cell Shape
- ...



# Types of Learning

- Supervised (inductive) learning
  - Given: training data + desired outputs (labels)
- Unsupervised learning
  - Given: training data (without desired outputs)
- Semi-supervised learning
  - Given: training data + a few desired outputs
- Reinforcement learning
  - Rewards from sequence of actions



**We will cover  
later in the class**

# 3. Decision Trees

# A learning problem: predict fuel efficiency

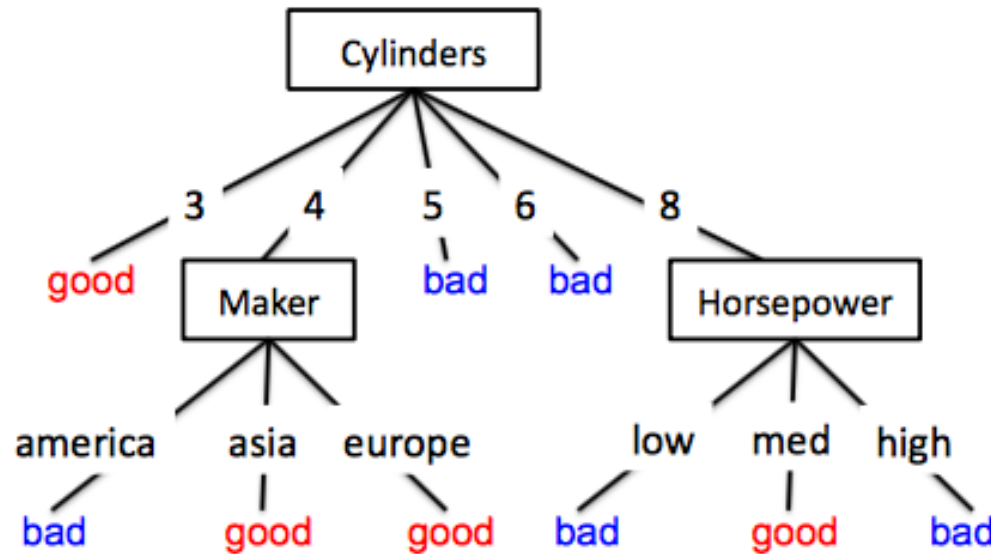
- 40 data points
- Goal: predict MPG
- Need to find:  
 $f : X \rightarrow Y$
- Discrete data (for now)

mpg	cylinders	displacement	horsepower	weight	acceleration	modelyear	maker
good	4	low	low	low	high	75to78	asia
bad	6	medium	medium	medium	medium	70to74	america
bad	4	medium	medium	medium	low	75to78	europa
bad	8	high	high	high	low	70to74	america
bad	6	medium	medium	medium	medium	70to74	america
bad	4	low	medium	low	medium	70to74	asia
bad	4	low	medium	low	low	70to74	asia
bad	8	high	high	high	low	75to78	america
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
bad	8	high	high	high	low	70to74	america
good	8	high	medium	high	high	79to83	america
bad	8	high	high	high	low	75to78	america
good	4	low	low	low	low	79to83	america
bad	6	medium	medium	medium	high	75to78	america
good	4	medium	low	low	low	79to83	america
good	4	low	low	medium	high	79to83	america
bad	8	high	high	high	low	70to74	america
good	4	low	medium	low	medium	75to78	europa
bad	5	medium	medium	medium	medium	75to78	europa

$Y$   $X$

# Hypotheses: decision trees $f: X \rightarrow Y$

- Each internal node tests an attribute  $x_i$
- Each branch assigns an attribute value  $x_i=v$
- Each leaf assigns a class  $y$
- To classify input  $x$ : traverse the tree from root to leaf, output the labeled  $y$




Human interpretable!

## Informal

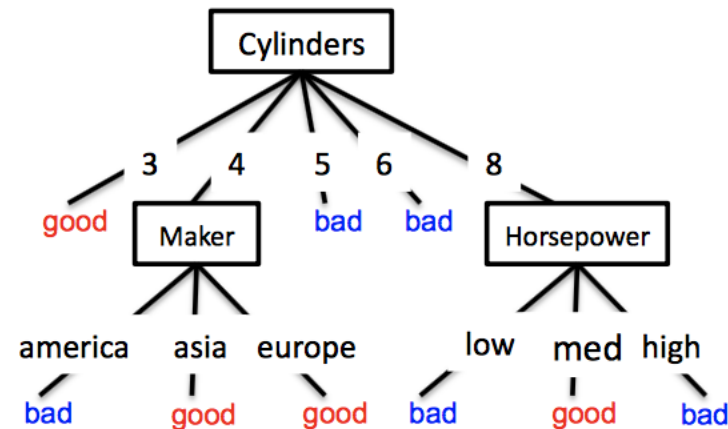
A hypothesis is a certain function that we believe (or hope) is similar to the true function, the *target function* that we want to model.

# What functions can Decision Trees represent?

- Decision trees can represent any function of the input attributes!
- For Boolean functions, path to leaf gives truth table row
- But, could require exponentially many nodes...



(Figure from Stuart Russell)

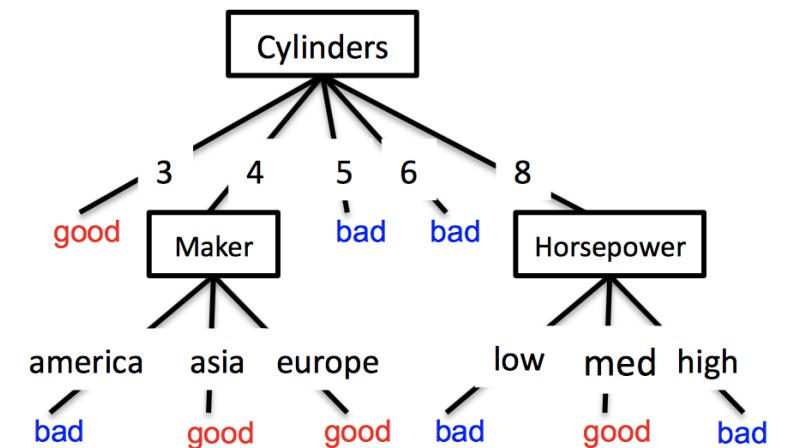


$\text{cyl}=3 \vee (\text{cyl}=4 \wedge (\text{maker}=\text{asia} \vee \text{maker}=\text{europe})) \vee \dots$

# Space of possible decision trees

- How will we choose the best one?
- Lets first look at how to split nodes, then consider how to find the best tree

mpg	cylinders	displacement	horsepower	weight	acceleration	modelyear	maker
good	4	low	low	low	high	75to78	asia
bad	6	medium	medium	medium	medium	70to74	america
bad	4	medium	medium	medium	low	75to78	europa
bad	8	high	high	high	low	70to74	america
bad	6	medium	medium	medium	medium	70to74	america
bad	4	low	medium	low	medium	70to74	asia
bad	4	low	medium	low	low	70to74	asia
bad	8	high	high	high	low	75to78	america
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
bad	8	high	high	high	low	70to74	america
good	8	high	medium	high	high	79to83	america
bad	8	high	high	high	low	75to78	america
good	4	low	low	low	low	79to83	america
bad	6	medium	medium	medium	high	75to78	america
good	4	medium	low	low	low	79to83	america
good	4	low	low	medium	high	79to83	america
bad	8	high	high	high	low	70to74	america
good	4	low	medium	low	medium	75to78	europa
bad	5	medium	medium	medium	medium	75to78	europa

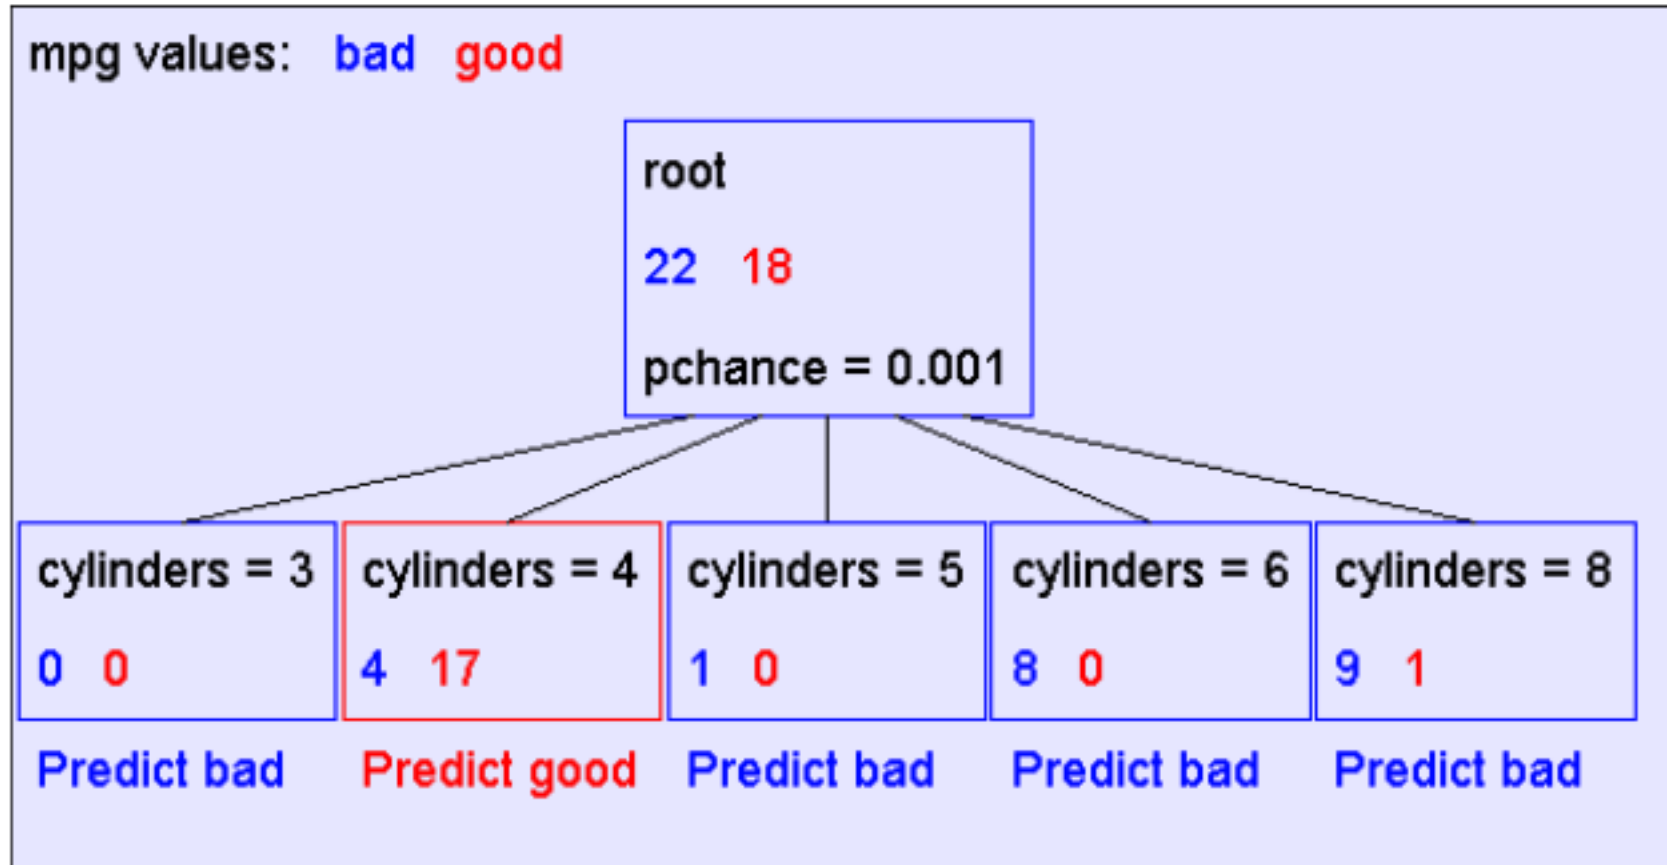


# What is the simplest tree?

- Always predict mpg = bad
  - We just take the majority class
- Is this a good tree?
  - We need to evaluate its performance
- Performance: We are correct on 22 examples and incorrect on 18 examples

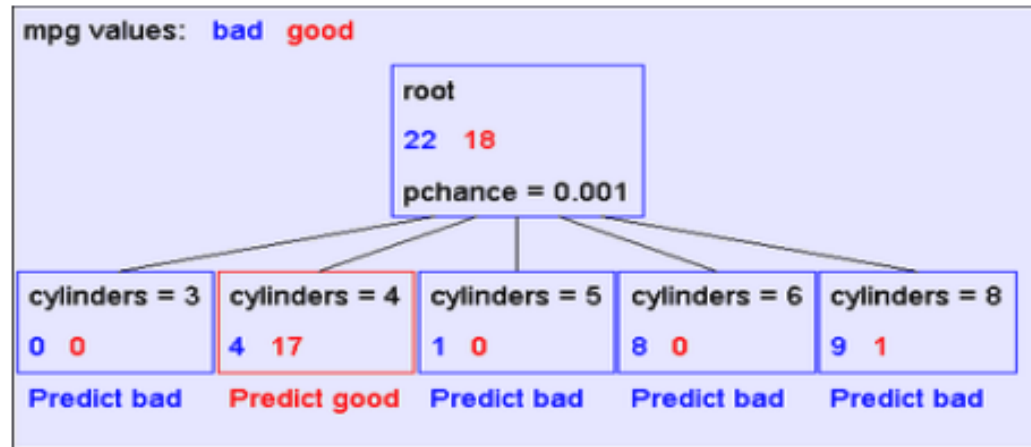
mpg	cylinders	displacement	horsepower	weight	acceleration	modelyear	maker
good	4	low	low	low	high	75to78	asia
bad	6	medium	medium	medium	medium	70to74	america
bad	4	medium	medium	medium	low	75to78	europa
bad	8	high	high	high	low	70to74	america
bad	6	medium	medium	medium	medium	70to74	america
bad	4	low	medium	low	medium	70to74	asia
bad	4	low	medium	low	low	70to74	asia
bad	8	high	high	high	low	75to78	america
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
:	:	:	:	:	:	:	:
bad	8	high	high	high	low	70to74	america
good	8	high	medium	high	high	79to83	america
bad	8	high	high	high	low	75to78	america
good	4	low	low	low	low	79to83	america
bad	6	medium	medium	medium	high	75to78	america
good	4	medium	low	low	low	79to83	america
good	4	low	low	medium	high	79to83	america
bad	8	high	high	high	low	70to74	america
good	4	low	medium	low	medium	75to78	europa
bad	5	medium	medium	medium	medium	75to78	europa

# A decision stump







# Recursive step



Take the  
Original  
Dataset..



And partition it  
according  
to the value of  
the attribute we  
split on



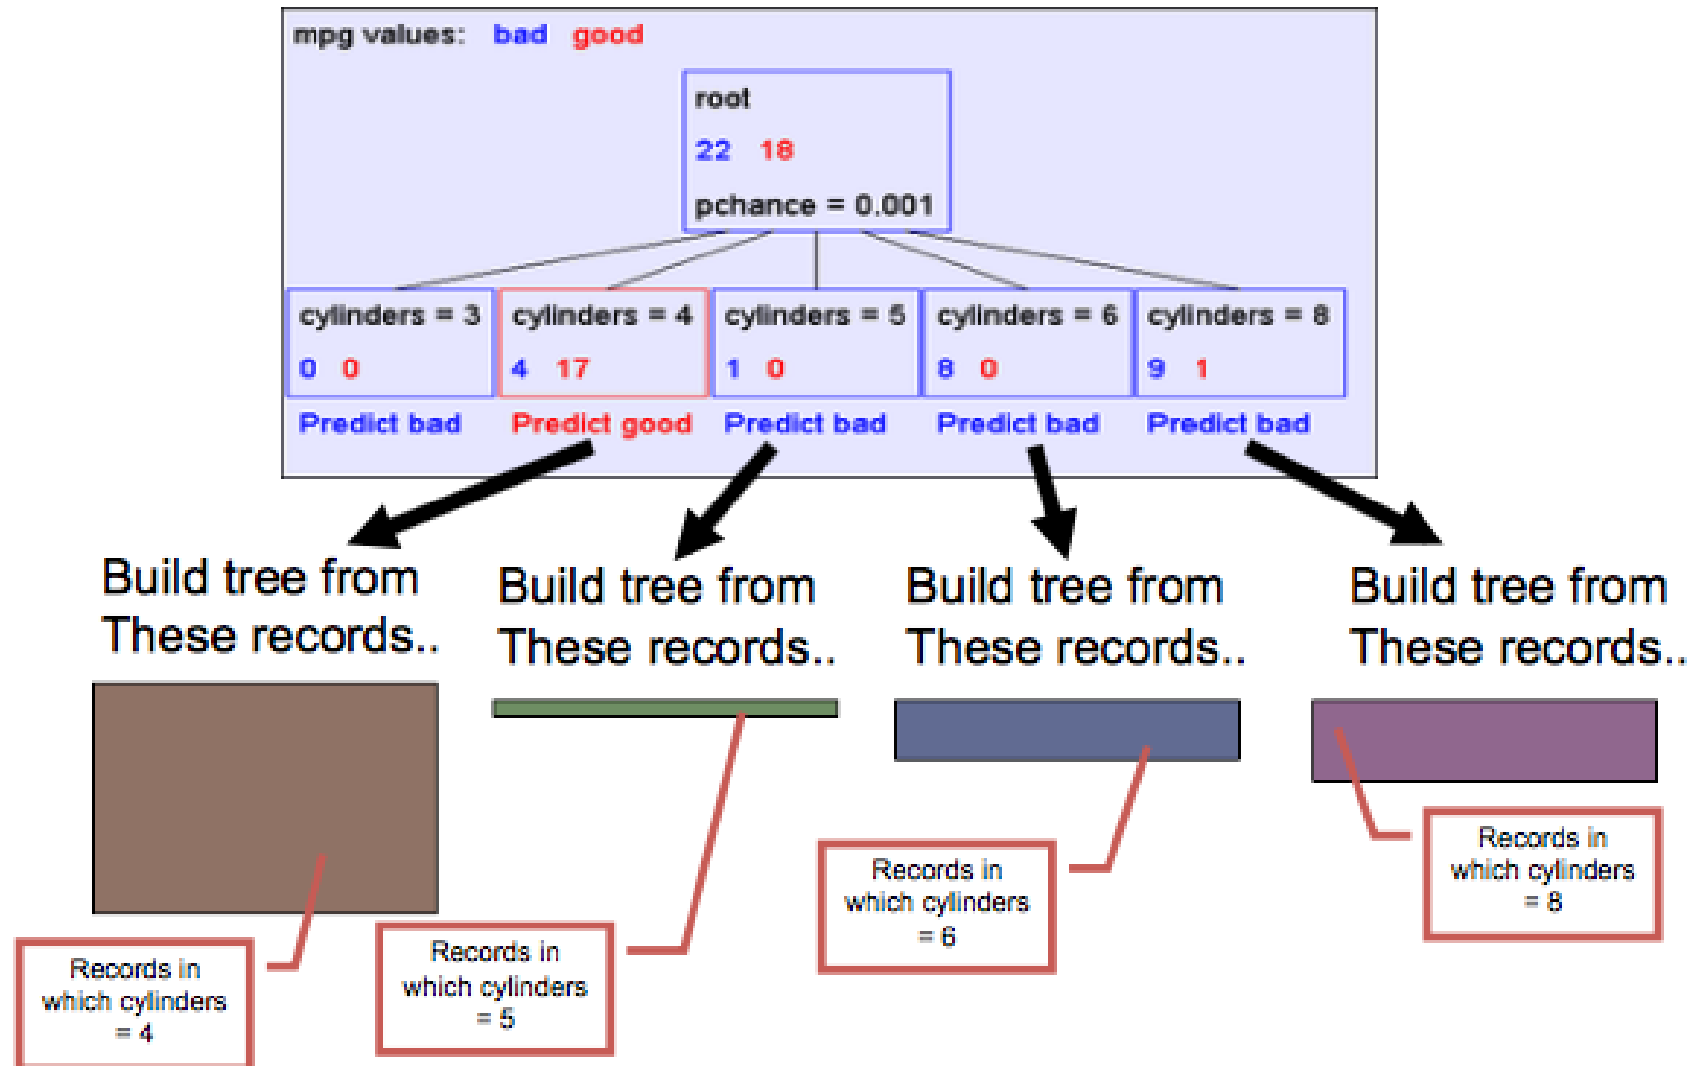
Records  
in which  
cylinders  
= 4

Records  
in which  
cylinders  
= 5

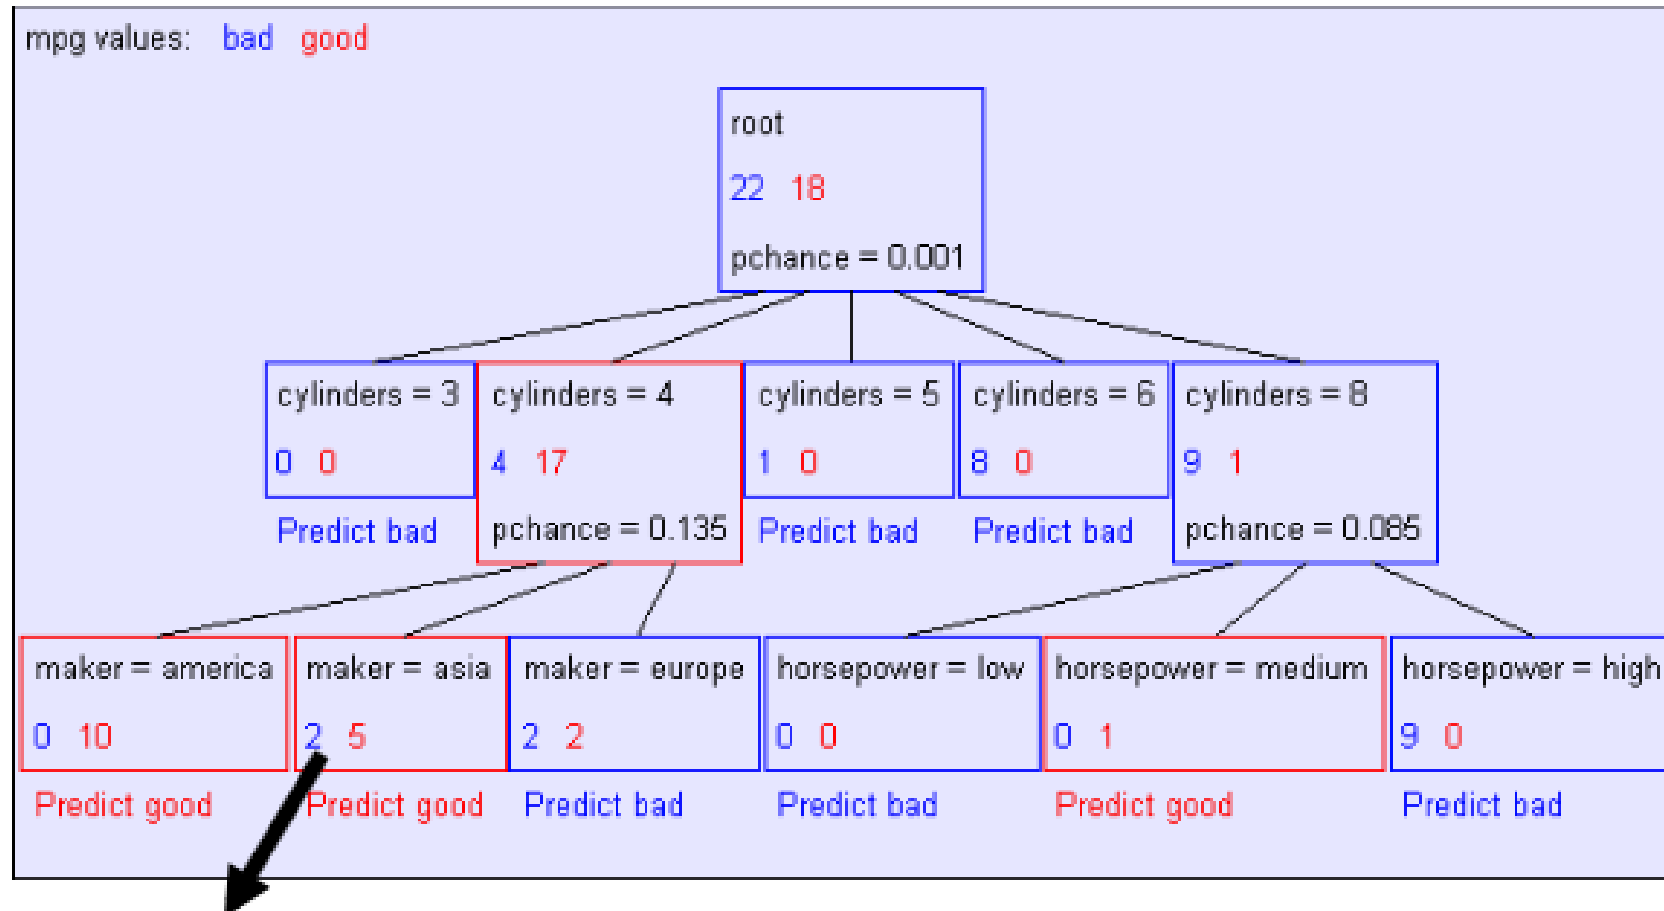
Records  
in which  
cylinders  
= 6

Records  
in which  
cylinders  
= 8

# Recursive step



# Second level of tree

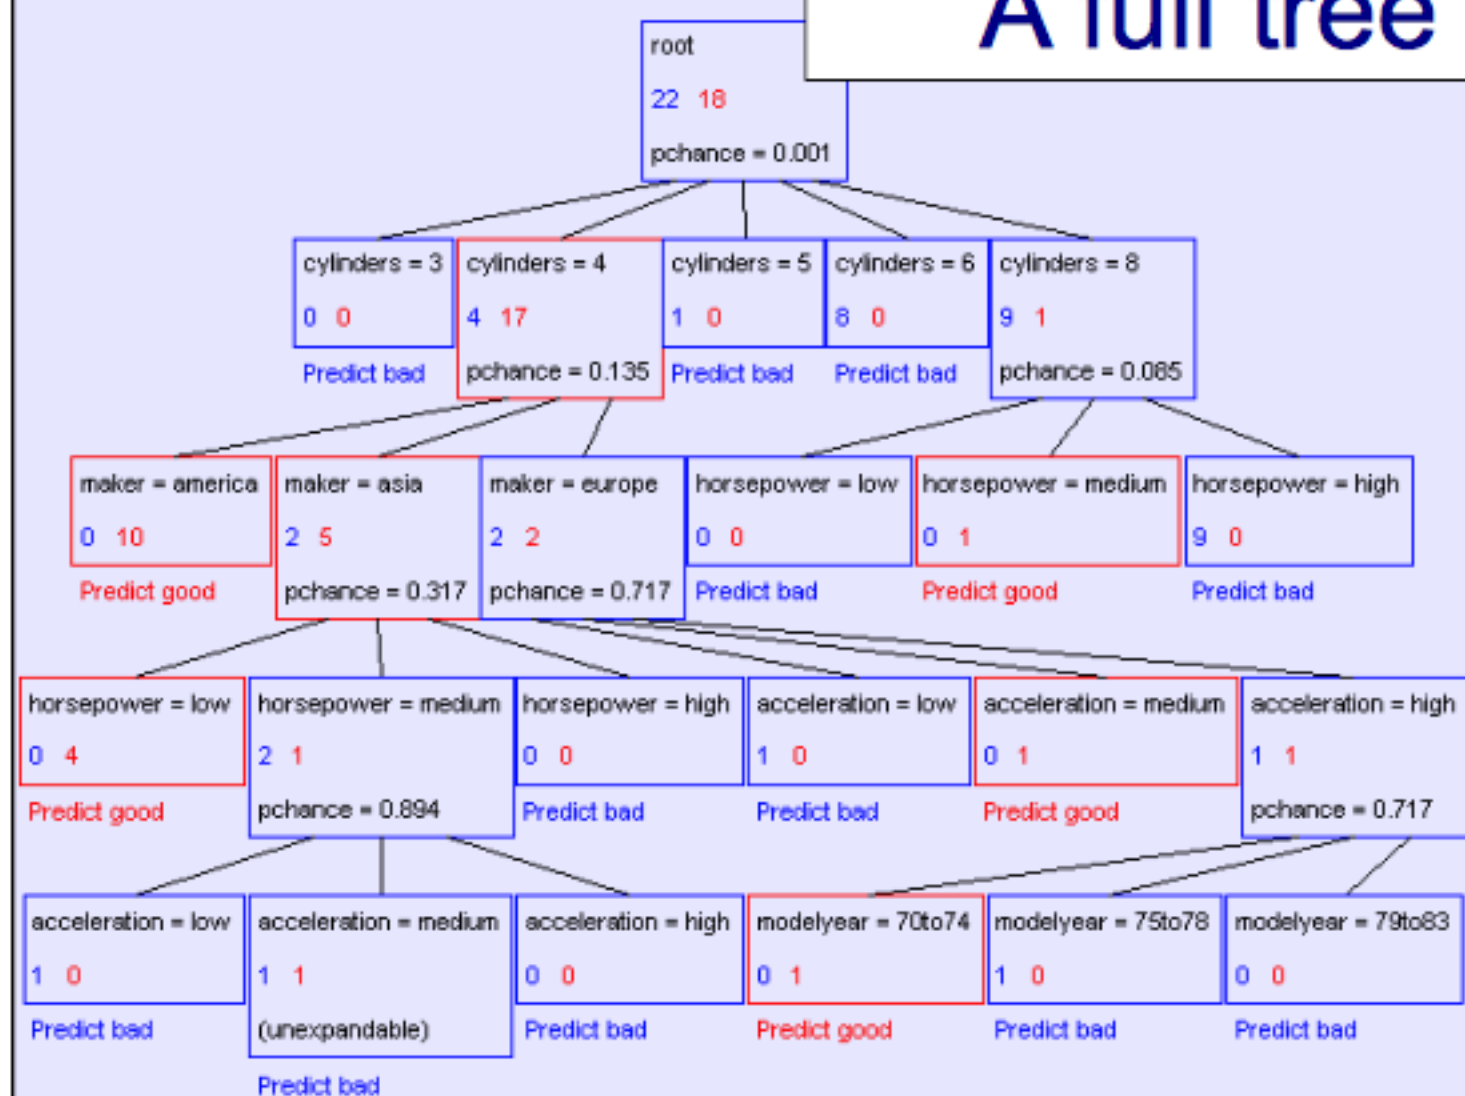


Recursively build a tree from the seven records in which there are four cylinders and the maker was based in Asia

(Similar recursion in the other cases)



# A full tree

mpg values: bad good



# Are all decision trees equal?

- Many trees can represent the same concept
- But, not all trees will have the same size!
  - e.g.,  $\phi = (A \wedge B) \vee (\neg A \wedge C) \dashv\vdash ((A \text{ and } B) \text{ or } (\text{not } A \text{ and } C))$




- Which tree do we prefer?

# Learning decision trees is hard

- Learning the simplest (smallest) decision tree is an NP-complete problem [Hyafil & Rivest '76]
- Resort to a greedy heuristic:
  - Start from empty decision tree
  - Split on **next best attribute (feature)**
  - Recurse

# Splitting: choosing a good attribute

Would we prefer to split on  $X_1$  or  $X_2$ ?



**Idea:** use counts at leaves to define probability distributions, so we can measure uncertainty!

$X_1$	$X_2$	$Y$
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F
F	T	F
F	F	F

# Measuring uncertainty

- Good split if we are more certain about classification after split
  - Deterministic good (all true or all false)
  - Uniform distribution bad
  - What about distributions in between?

$P(Y=A) = 1/2$	$P(Y=B) = 1/4$	$P(Y=C) = 1/8$	$P(Y=D) = 1/8$
----------------	----------------	----------------	----------------

$P(Y=A) = 1/4$	$P(Y=B) = 1/4$	$P(Y=C) = 1/4$	$P(Y=D) = 1/4$
----------------	----------------	----------------	----------------



# Entropy


Entropy  $H(Y)$  of a random variable  $Y$

$$H(Y) = - \sum_{i=1}^k P(Y = y_i) \log_2 P(Y = y_i)$$

**More uncertainty, more entropy!**

*Information Theory interpretation:*

$H(Y)$  is the expected number of bits needed to encode a randomly drawn value of  $Y$  (under most efficient code)



# High, Low Entropy

- “High Entropy”
  - Y is from a uniform like distribution
  - Flat histogram
  - Values sampled from it are less predictable
- “Low Entropy”
  - Y is from a varied (peaks and valleys) distribution
  - Histogram has many lows and highs
  - Values sampled from it are more predictable


# Entropy Example

$$H(Y) = - \sum_{i=1}^k P(Y = y_i) \log_2 P(Y = y_i)$$

$$P(Y=t) = 5/6$$

$$P(Y=f) = 1/6$$

$$\begin{aligned} H(Y) &= - 5/6 \log_2 5/6 - 1/6 \log_2 1/6 \\ &= 0.65 \end{aligned}$$



$X_1$	$X_2$	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F

# Conditional Entropy

Conditional Entropy  $H(Y|X)$  of a random variable  $Y$  conditioned on a random variable  $X$


$$H(Y | X) = - \sum_{j=1}^v P(X = x_j) \sum_{i=1}^k P(Y = y_i | X = x_j) \log_2 P(Y = y_i | X = x_j)$$

Example:

$$P(X_1 = \text{t}) = 4/6$$

$$P(X_1 = \text{f}) = 2/6$$

$$\begin{aligned} H(Y|X_1) &= - 4/6 (1 \log_2 1 + 0 \log_2 0) \\ &\quad - 2/6 (1/2 \log_2 1/2 + 1/2 \log_2 1/2) \\ &= 2/6 \end{aligned}$$



$X_1$	$X_2$	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F

# Information gain

- Decrease in entropy (uncertainty) after splitting

$$IG(X) = H(Y) - H(Y | X)$$

In our running example:

$$\begin{aligned} IG(X_1) &= H(Y) - H(Y|X_1) \\ &= 0.65 - 0.33 \end{aligned}$$

$IG(X_1) > 0 \rightarrow$  we prefer the split!

$X_1$	$X_2$	Y
T	T	T
T	F	T
T	T	T
T	F	T
F	T	T
F	F	F

# Learning decision trees


- Start from empty decision tree
- Split on **next best attribute (feature)**
  - Use, for example, information gain to select attribute:

$$\arg \max_i IG(X_i) = \arg \max_i H(Y) - H(Y | X_i)$$

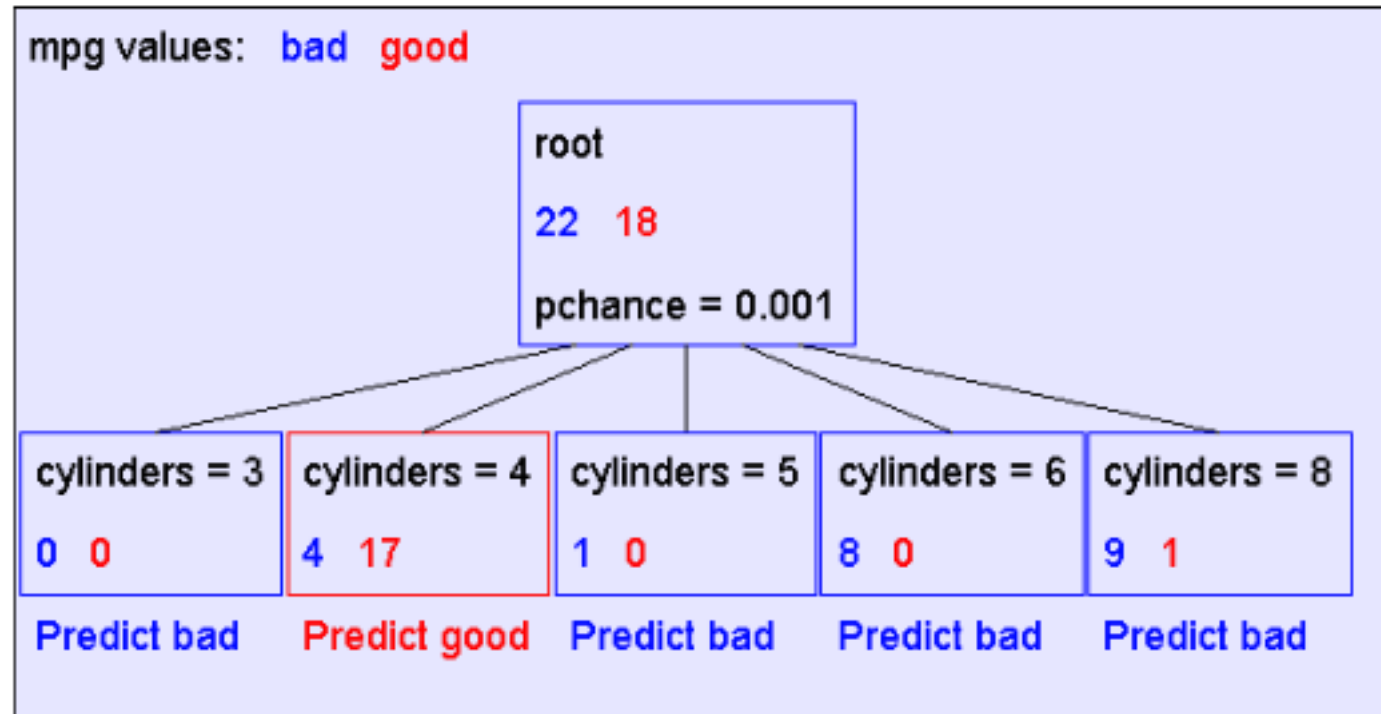
- Recurse

Suppose we want  
to predict MPG

Look at all the  
information  
gains...



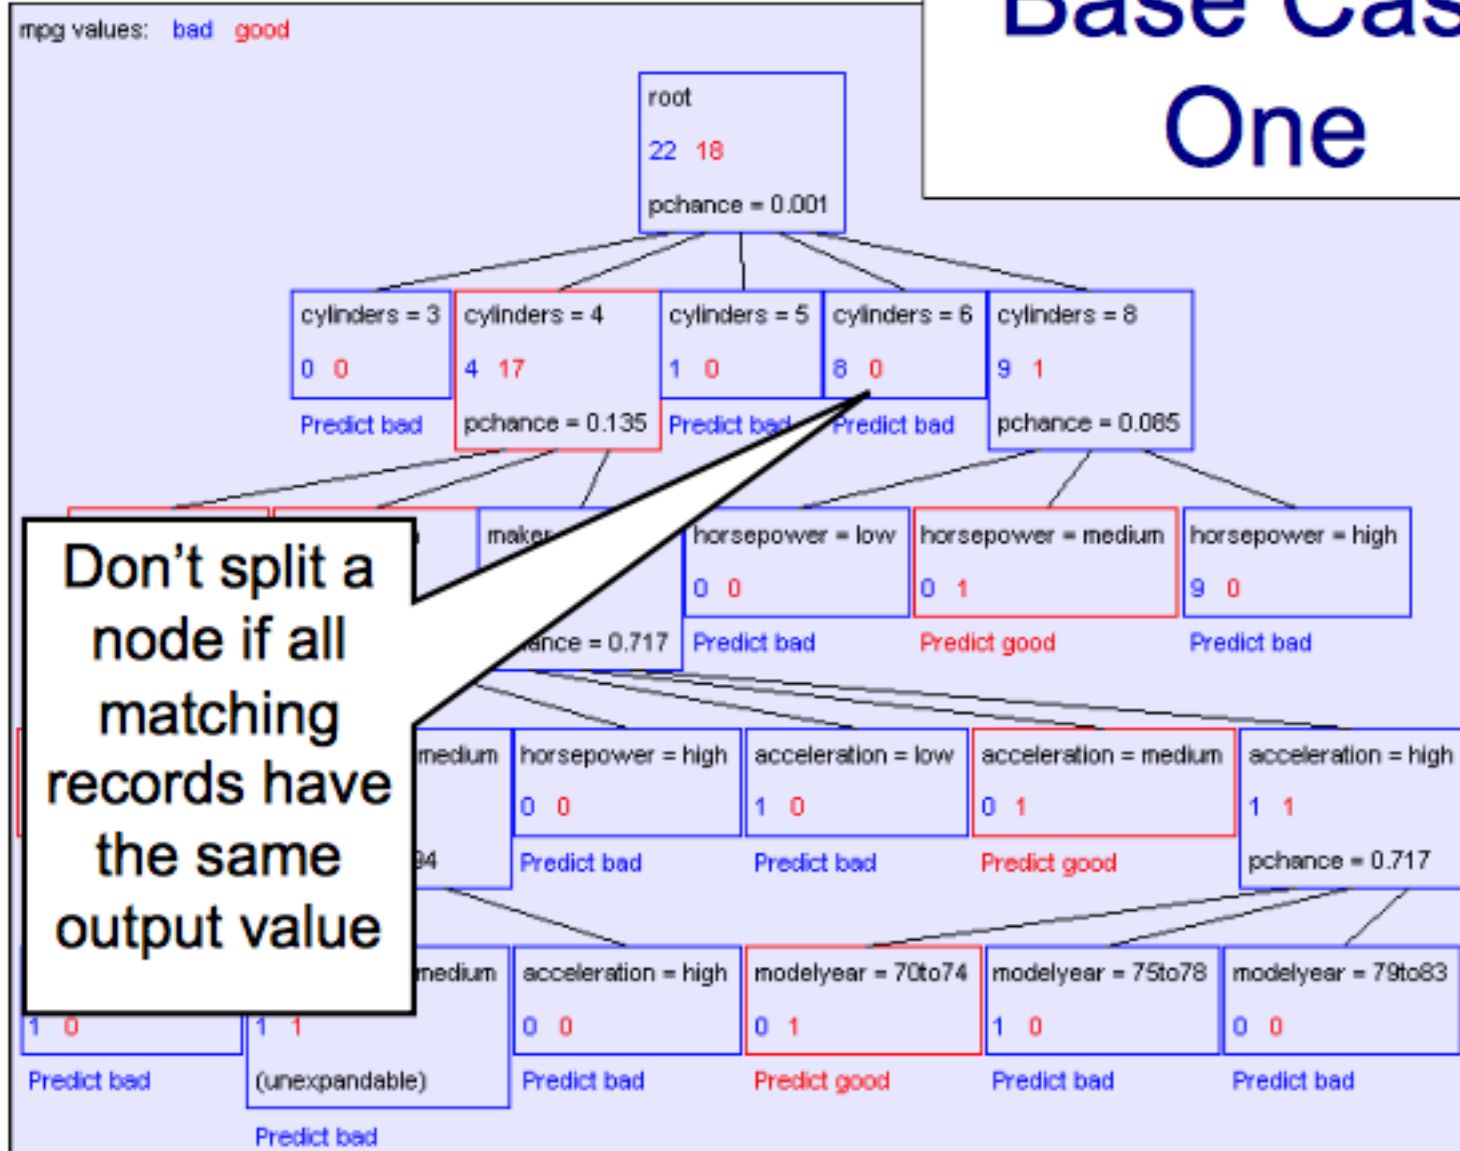
# A decision stump



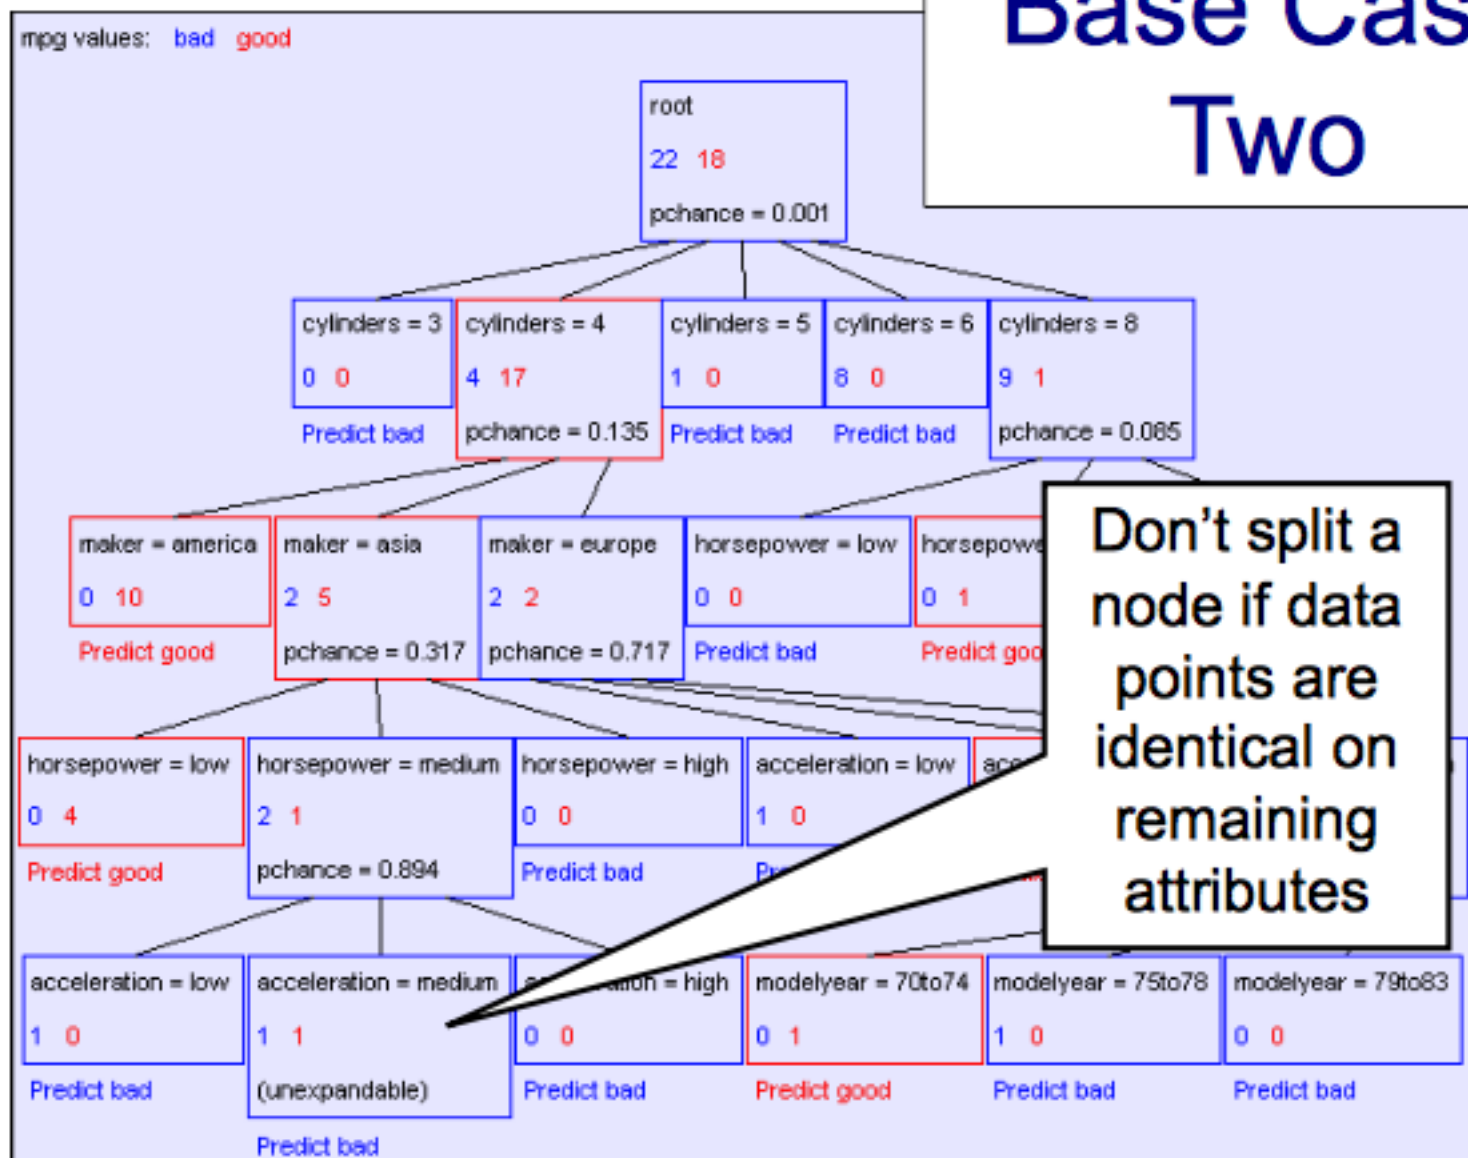
First split looks good! But, when do we stop?




# Base Case One



# Base Case Two



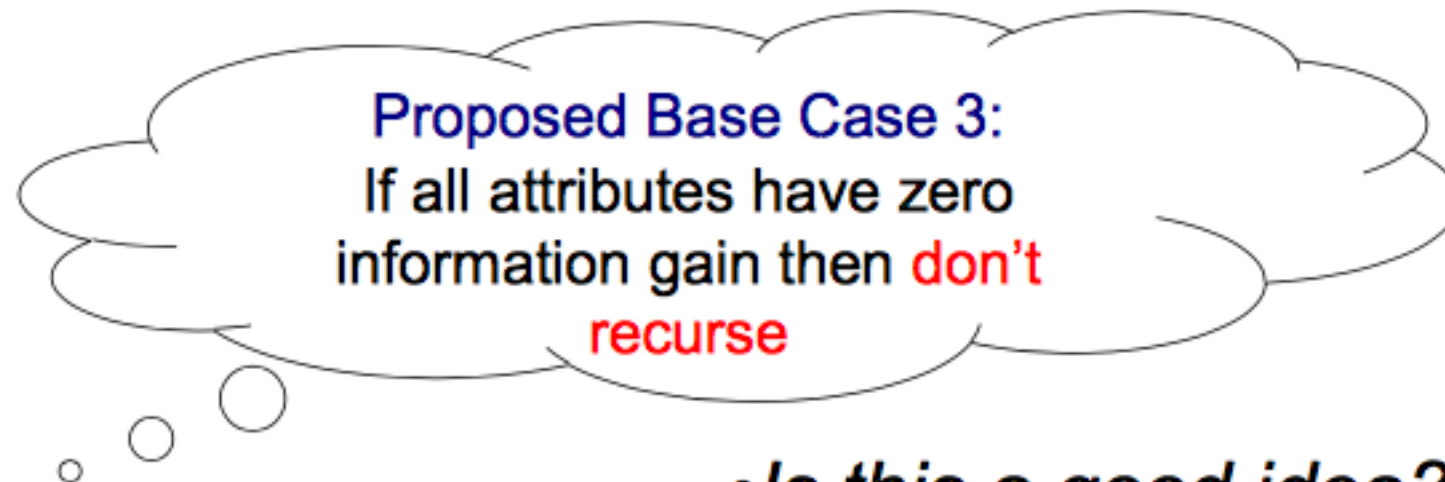
## Base Case Two: No attributes can distinguish



Information gains using the training set (2 records)			
mpg values: bad good			
Input	Value	Distribution	Info Gain
cylinders	3		0
	4		
	5		
	6		
	8		
displacement	low		0
	medium		
	high		
horsepower	low		0
	medium		
	high		
weight	low		0
	medium		
	high		
acceleration	low		0
	medium		
	high		
modelyear	70to74		0
	75to78		
	79to83		
maker	america		0
	asia		
	europa		

# Base Cases: An Idea

- **Base Case One:** If all records in current data subset have the same output then **do not recurse**
- **Base Case Two:** If all records have exactly the same set of input attributes then **do not recurse**



• *Is this a good idea?*

# The problem with Base Case 3


$$y = a \text{ XOR } b$$

a	b	y
0	0	0
0	1	1
1	0	1
1	1	0

The information gains:

Information gains using the training set (4 records)				
y values: 0 1				
Input	Value	Distribution	Info Gain	
a	0		0	
	1			
b	0		0	
	1			

The resulting decision tree:




# If we omit Base Case 3

$$y = a \text{ XOR } b$$

a	b	y
0	0	0
0	1	1
1	0	1
1	1	0

Is it OK to omit Base Case 3?

The resulting decision tree:



# Summary: Building Decision Trees

**BuildTree(*DataSet*, *Output*)**


- If all output values are the same in *DataSet*, return a leaf node that says “predict this unique output”
- If all input values are the same, return a leaf node that says “predict the majority output”
- Else find attribute *X* with highest Info Gain
- Suppose *X* has  $n_X$  distinct values (i.e. *X* has arity  $n_X$ ).
  - Create a non-leaf node with  $n_X$  children.
  - The *i*’th child should be built by calling

**BuildTree(*DS<sub>i</sub>*, *Output*)**

Where *DS<sub>i</sub>* contains the records in *DataSet* where *X* = *i*th value of *X*.

# From categorical to real-valued attributes

- **Binary tree:** split on attribute  $X$  at value  $t$ 
  - One branch:  $X < t$
  - Other branch:  $X \geq t$
- **Requires small change**
  - Allow repeated splits on same variable
  - How does this compare to “branch on each value” approach?





# What you need to know about decision trees

- Decision trees are one of the most popular ML tools
  - Easy to understand, implement, and use
  - Computationally cheap (to solve heuristically)
- Information gain to select attributes
- Presented for classification but can be used for regression and density estimation too
- Decision trees will overfit!!!
  - We will see the definition of overfitting and related concepts later in class.