

# Data Management for Data Science

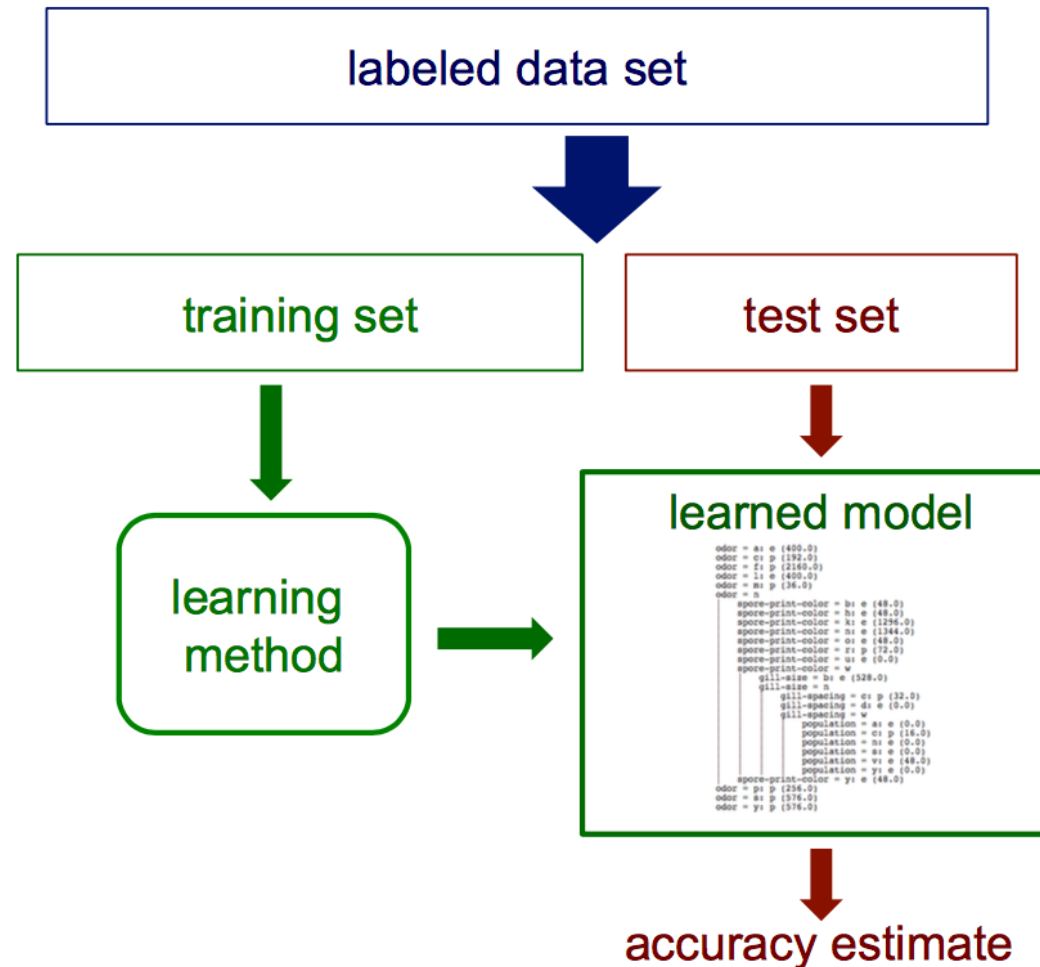
Lecture 19: Evaluating Machine Learning Methods

Prof. Asoc. Endri Raço

# Today

1. Evaluating ML models

# How can we get an unbiased estimate of the accuracy of a learned model?




# Test sets

- How can we get an unbiased estimate of the accuracy of a learned model?
- When learning a model, you should pretend that you don't have the test data yet
- If the test-set labels influence the learned model in any way, accuracy estimates will be biased


# Learning curves

- How does the accuracy of a learning method change as a function of the training-set size?
  - This can be assessed by learning curves



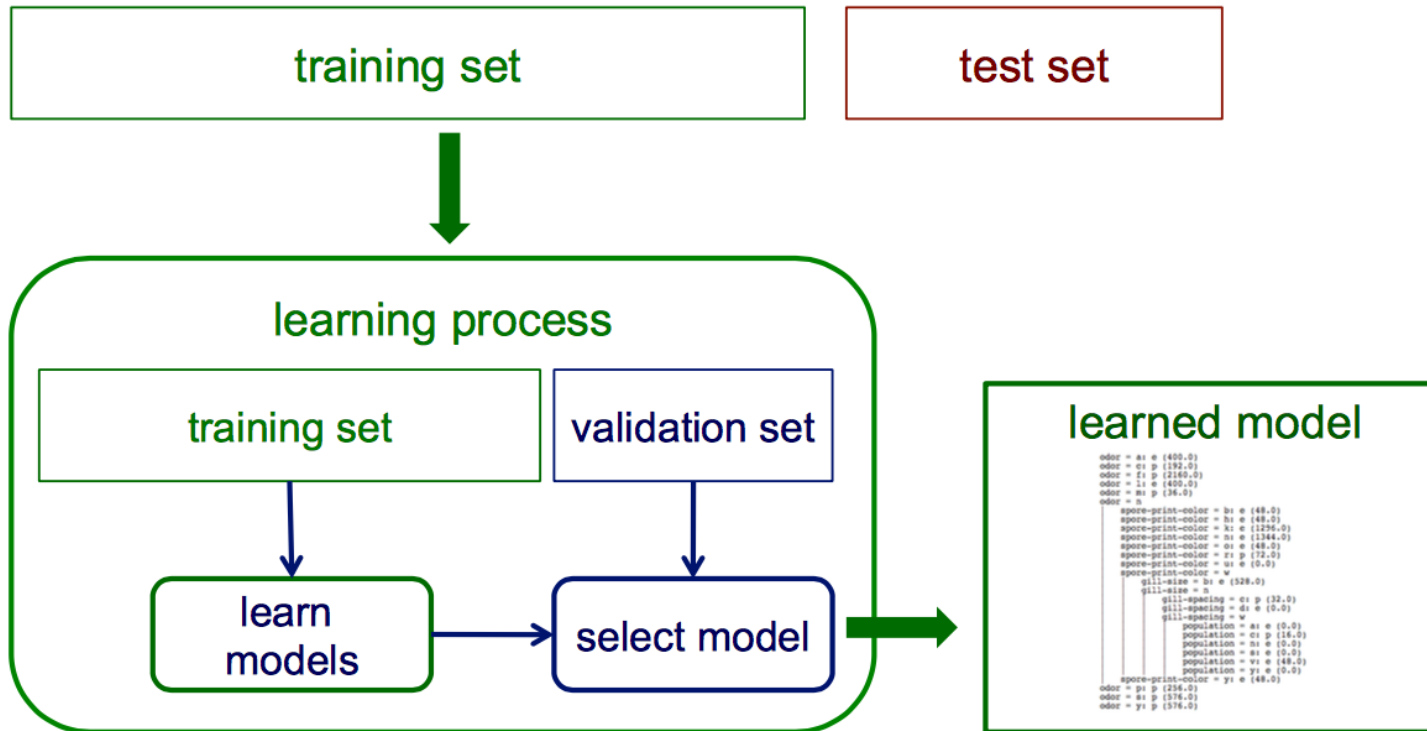
# Learning curves

- Given a training/test set partition
  - For each sample size  $s$  on the learning curve
    - (optionally) repeat  $n$  times
    - Randomly select  $s$  instances from the training set
    - Learn the model
    - Evaluate the model on the test set to determine accuracy  $a$
    - Plot  $(s, a)$



# Validation (tuning) sets

- Suppose we want unbiased estimates of accuracy during the learning process (e.g. to choose the best level of decision-tree pruning)?



Partition training data into separate training/validation sets

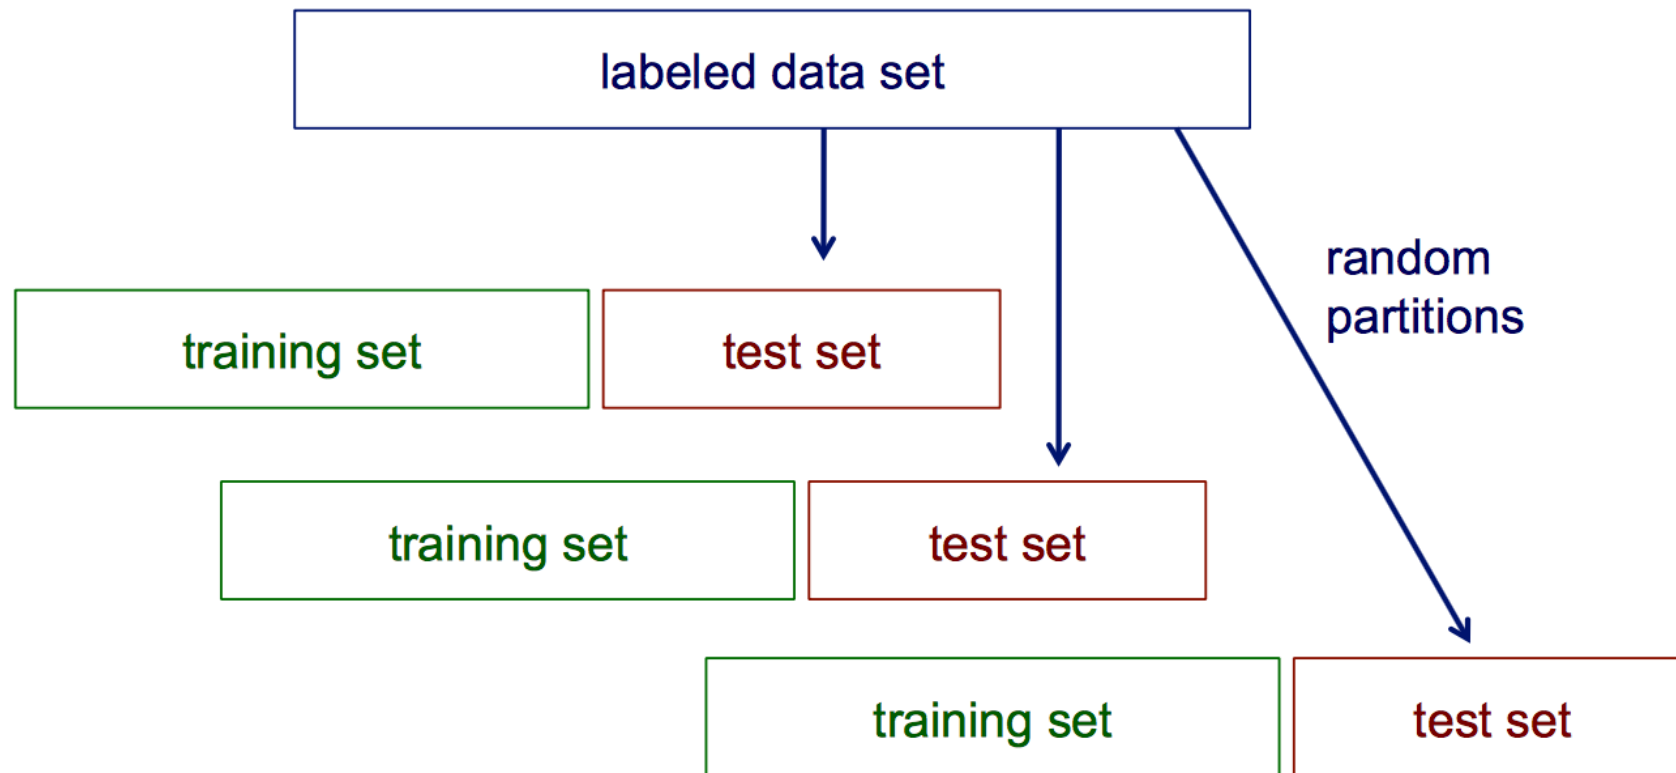
# Limitations of using a single training/test partition

- We may not have enough data to make sufficiently large training and test sets
  - A **larger test set** gives us more reliable estimates of accuracy (i.e., a lower variance estimate)
  - But... a **larger training set** will be more representative of how much data we actually have for learning process
- A single training set does not tell us how sensitive accuracy is to a particular training sample



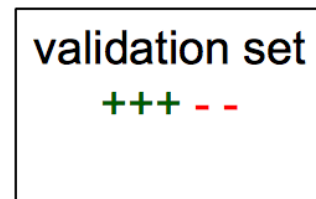
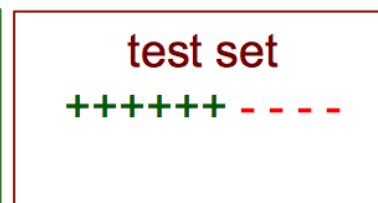
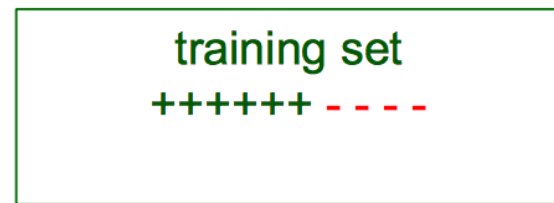
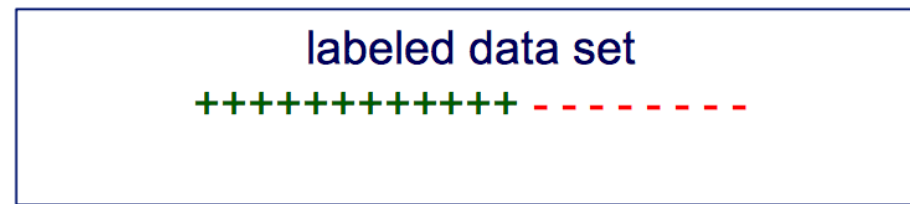
# Random resampling

- We can address the second issue by repeatedly randomly partitioning the available data into training and set sets.



# Stratified sampling


- When randomly selecting training or validation sets, we may want to ensure that class proportions are maintained in each selected set



This can be done via stratified sampling: first stratify instances by class, then randomly select instances from each class proportionally.

# Cross validation

partition data  
into  $n$  subsamples



iteratively leave one  
subsample out for  
the test set, train on  
the rest

iteration	train on	test on
1	$s_2$ $s_3$ $s_4$ $s_5$	$s_1$
2	$s_1$ $s_3$ $s_4$ $s_5$	$s_2$
3	$s_1$ $s_2$ $s_4$ $s_5$	$s_3$
4	$s_1$ $s_2$ $s_3$ $s_5$	$s_4$
5	$s_1$ $s_2$ $s_3$ $s_4$	$s_5$

# Cross validation example

- Suppose we have 100 instances, and we want to estimate accuracy with cross validation

iteration	train on	test on	correct
1	$s_2$ $s_3$ $s_4$ $s_5$	$s_1$	11 / 20
2	$s_1$ $s_3$ $s_4$ $s_5$	$s_2$	17 / 20
3	$s_1$ $s_2$ $s_4$ $s_5$	$s_3$	16 / 20
4	$s_1$ $s_2$ $s_3$ $s_5$	$s_4$	13 / 20
5	$s_1$ $s_2$ $s_3$ $s_4$	$s_5$	16 / 20

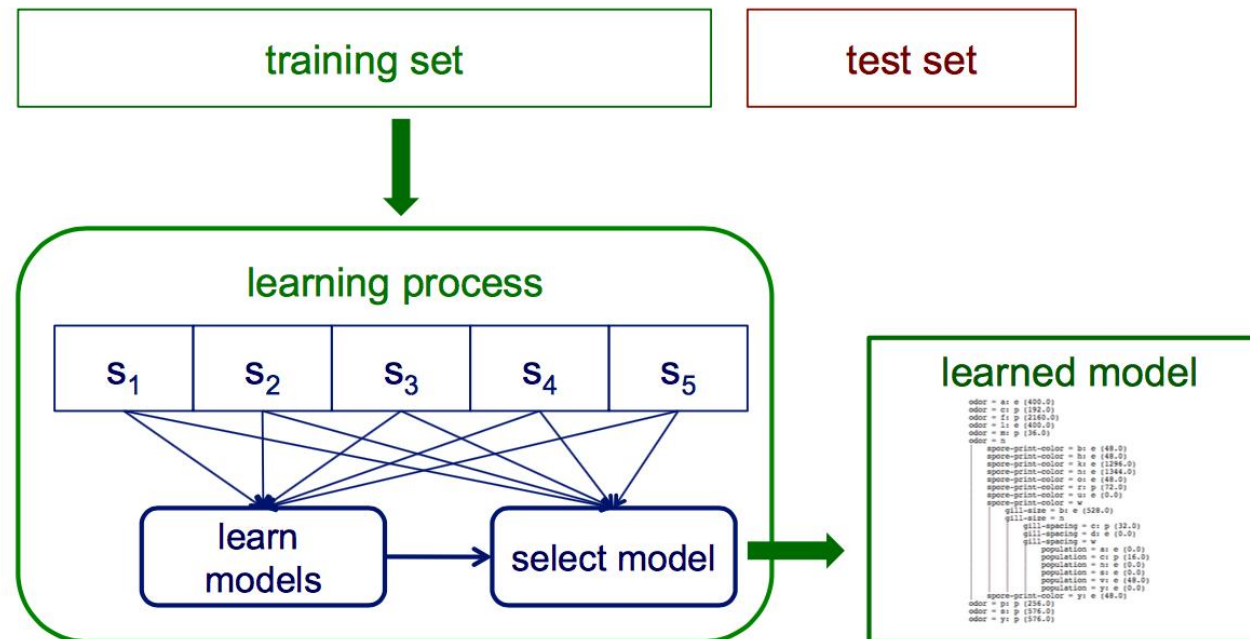
accuracy = 73/100 = 73%

# Cross validation example

- 10-fold cross validation is common, but smaller values of  $n$  are often used when learning takes a lot of time
- In *leave-one-out* cross validation,  $n = \text{\#instances}$
- In *stratified* cross validation, stratified sampling is used when partitioning the data
- CV makes efficient use of the available data for testing
- Note that whenever we use multiple training sets, as in CV and random resampling, we are evaluating a learning method as opposed to an individual learned model


# Internal cross validation

- Instead of a single validation set, we can use cross-validation within a training set to select a model (e.g. to choose the best level of decision-tree pruning)



# Confusion matrices

- How can we understand what types of mistakes a learned model makes?



# Confusion matrix for 2-class problems

		actual class	
		positive	negative
predicted class	positive	true positives (TP)	false positives (FP)
	negative	false negatives (FN)	true negatives (TN)

$$\text{accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{FP} + \text{FN} + \text{TN}}$$



# Is accuracy an adequate measure of predictive performance?

- accuracy may not be useful measure in cases where
  - there is a large class skew
    - Is 98% accuracy good if 97% of the instances are negative?
- there are differential misclassification costs – say, getting a positive wrong costs more than getting a negative wrong
  - Consider a medical domain in which a false positive results in an extraneous test but a false negative results in a failure to treat a disease
- we are most interested in a subset of high-confidence predictions

# Other accuracy metrics


		actual class	
		positive	negative
predicted class	positive	true positives (TP)	false positives (FP)
	negative	false negatives (FN)	true negatives (TN)

$$\text{true positive rate (recall)} = \frac{\text{TP}}{\text{actual pos}} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{false positive rate} = \frac{\text{FP}}{\text{actual neg}} = \frac{\text{FP}}{\text{TN} + \text{FP}}$$

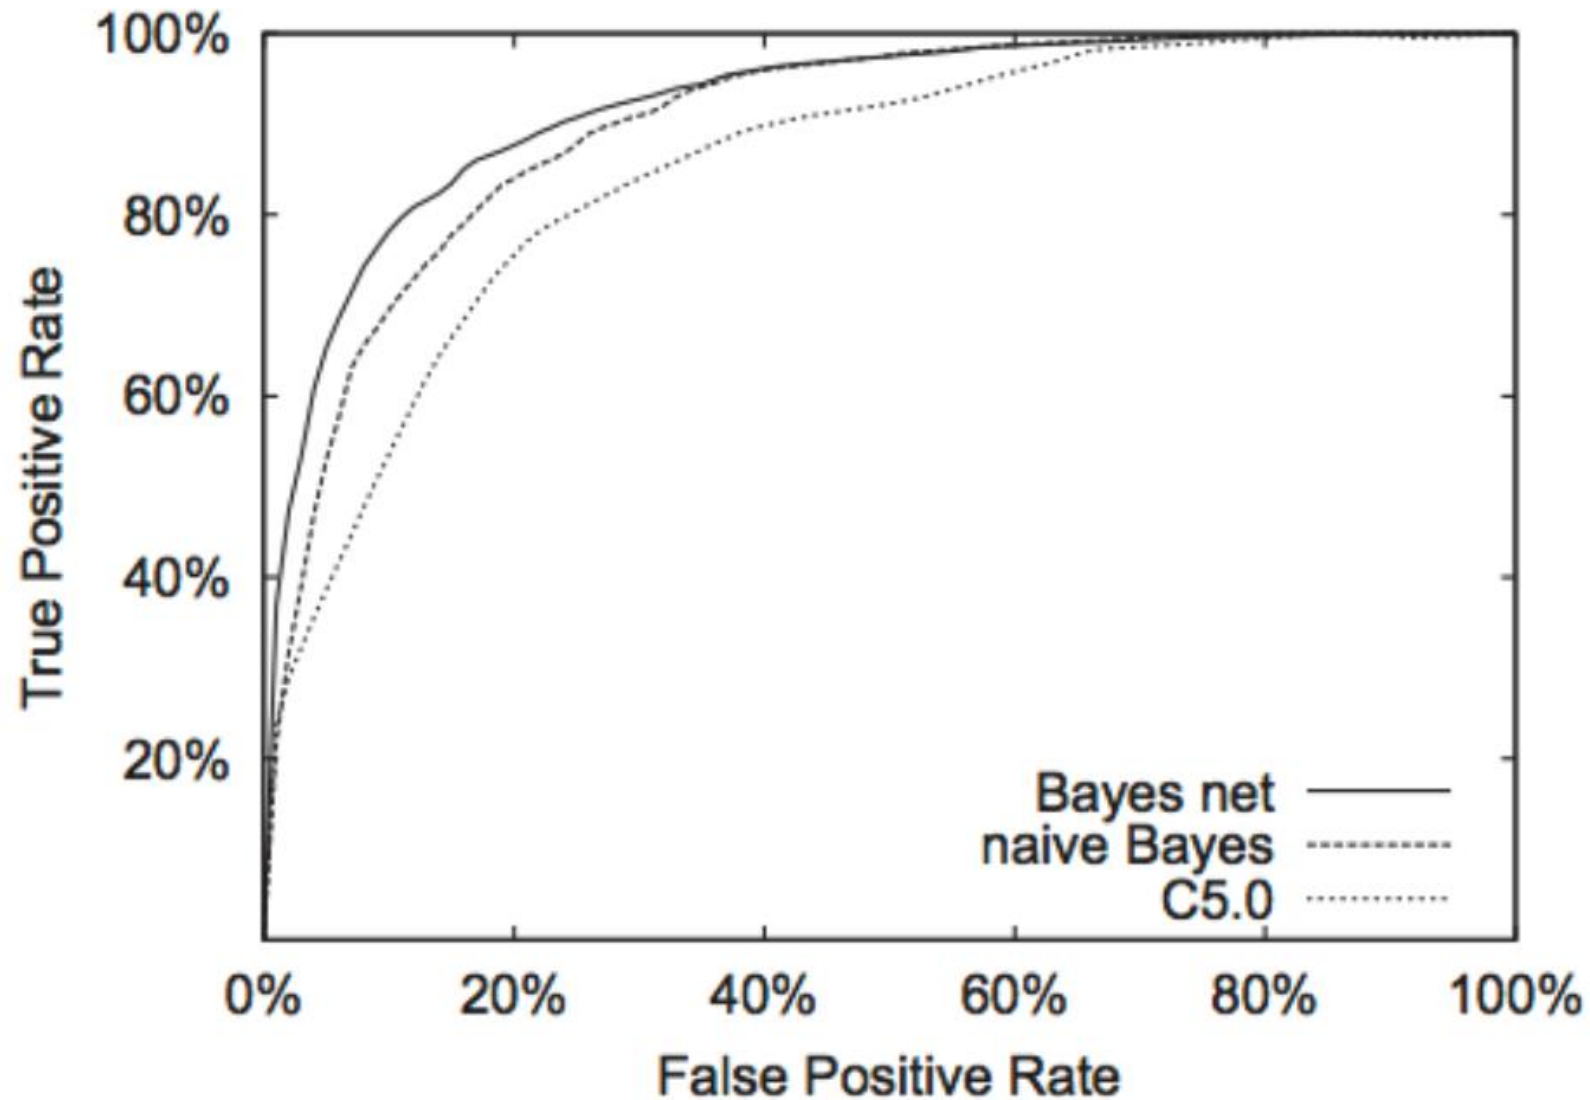
# ROC curves

A *Receiver Operating Characteristic (ROC)* curve plots the TP-rate vs. the FP-rate as a threshold on the confidence of an instance being positive is varied

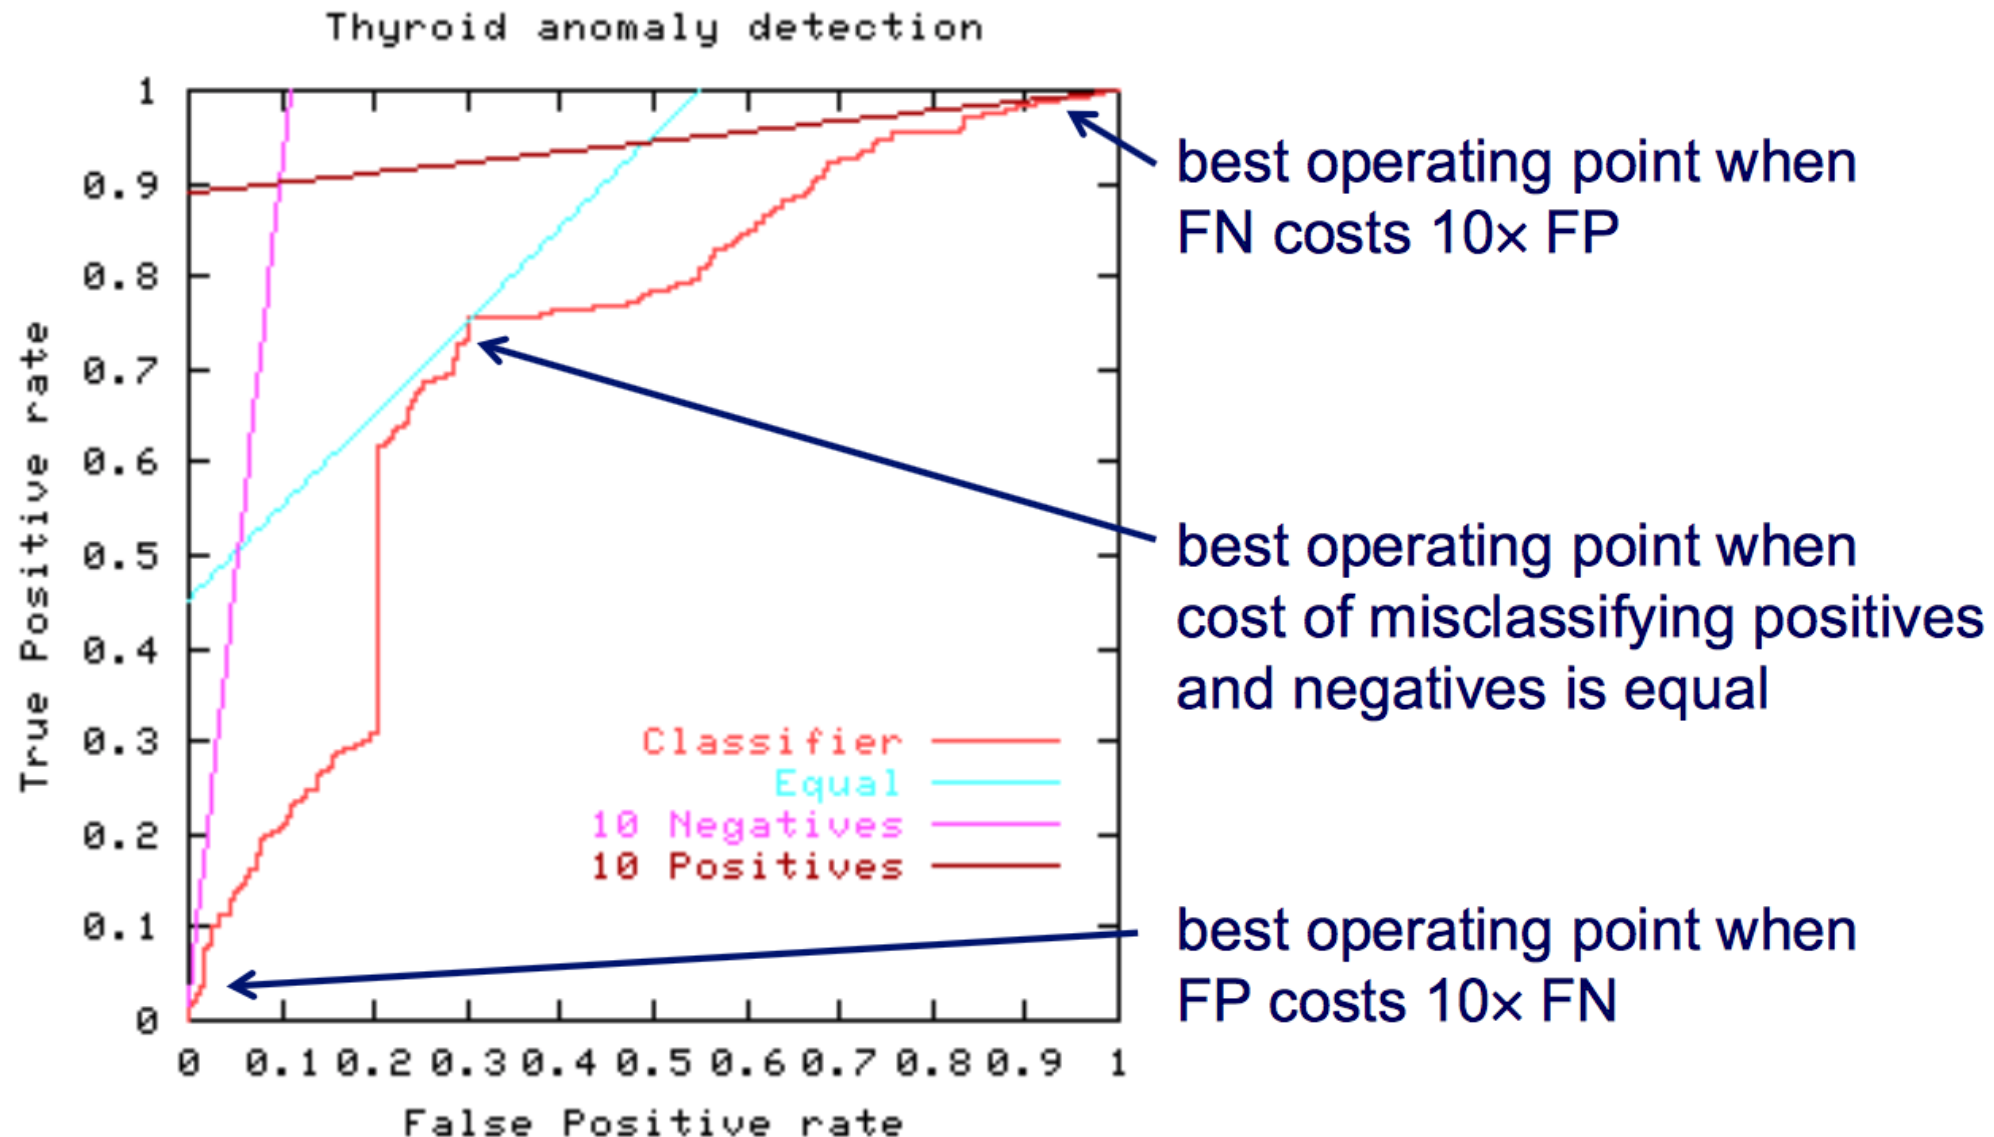


Different methods can work better in different parts of ROC space. This depends on cost of false + vs. false -

# ROC curve example



# ROC curves and misclassification costs



# Algorithm for creating an ROC curve

1. sort test-set predictions according to confidence that each instance is positive
2. step through sorted list from high to low confidence
  - i. locate a *threshold* between instances with opposite classes (keeping instances with the same confidence value on the same side of threshold)
  - ii. compute TPR, FPR for instances above threshold
  - iii. output (FPR, TPR) coordinate

# Other accuracy metrics


		actual class	
		positive	negative
predicted class	positive	true positives (TP)	false positives (FP)
	negative	false negatives (FN)	true negatives (TN)

$$\text{recall (TP rate)} = \frac{\text{TP}}{\text{actual pos}} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{precision} = \frac{\text{TP}}{\text{predicted pos}} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

# Precision/recall curves

*A precision/recall curve plots the precision vs. recall (TP-rate) as a threshold on the confidence of an instance being positive is varied*





# To Avoid Cross-Validation Pitfalls

- 1. Is my held-aside test data really representative of going out to collect new data?
  - Even if your methodology is fine, someone may have collected features for positive examples differently than for negatives – should be *randomized*
  - Example: samples from cancer processed by different people or on different days than samples for normal controls

# To Avoid Cross-Validation Pitfalls

- 2. Did I repeat my entire data processing procedure on every fold of cross-validation, using only the training data for that fold?
  - On each fold of cross-validation, did I ever access in any way the label of a test case?
  - Any preprocessing done over *entire data set* (feature selection, parameter tuning, threshold selection) must *not* use labels

# To Avoid Cross-Validation Pitfalls

- 3. Have I modified my algorithm so many times, or tried so many approaches, on this same data set that *I* (the human) am overfitting it?
  - Have I continually modified my preprocessing or learning algorithm until I got some improvement on this data set?
  - If so, I really need to get some additional data now to at least test on

# Ablation Studies

We can gain insight into what contributes to a learning system's performance by removing (lesioning) components of it

The ROC curves here show how performance is affected when various feature types are removed from the learning representation

